

AB-100: Architect AI Solutions for Business Productivity

80 Advanced Practice Questions — Answer Key with Rationale

Covers: Manage testing AI-powered business solutions · Design ALM process for AI-powered business solutions · Design responsible AI security, governance, risk management, and compliance.

Question types included: Single-answer MCQ, Multi-select MCQ, Drop-down selection, Drag-and-drop matching, and Hotspot/case-study Yes-No statement sets — matching the style of official AB-100 assessment items.

Section A — Manage Testing AI-Powered Business Solutions (Q1–Q12)

Q1. MULTIPLE CHOICE — SINGLE ANSWER

A solution architect is designing a validation approach for a new Copilot-based expense-approval agent. Which statement best justifies why traditional, deterministic test scripts alone are insufficient?

- A. Deterministic scripts cannot be automated in CI/CD pipelines.
- B. AI outputs are probabilistic and vary with context, data, and phrasing, so fixed expected-output assertions will produce false failures.
- C. Deterministic scripts require more storage than AI evaluation datasets.
- D. Dynamics 365 does not support scripted UI testing.

Correct Answer: B

Why: Generative AI responses are non-deterministic — identical inputs can yield different (but still valid) outputs depending on context, grounding data, and prompt phrasing. Traditional pass/fail scripts assume a single correct output and therefore misclassify valid variation as failure; AI solutions need scenario-based, threshold-driven evaluation instead.

Q2. MULTIPLE CHOICE — SINGLE ANSWER

You are asked to select the single most important metric for confirming that an AI solution's outputs support the intended business outcome. Which should you choose?

- A. Number of concurrent users
- B. Accuracy and relevance of the AI's responses
- C. Frequency of UI updates
- D. Length of the conversation

Correct Answer: B

Why: Business alignment is measured by whether responses are correct and relevant to the task — UI cosmetics, session length, and user counts are operational signals, not outcome-quality signals.

Q3. MULTIPLE CHOICE — SINGLE ANSWER

An AI solution spans Dynamics 365 Sales, Customer Service, and a Copilot Studio agent. Why must end-to-end tests validate cross-app data flow rather than testing each app in isolation?

- A. Each Dynamics 365 app requires a dedicated AI model that cannot be shared.
- B. AI output quality depends on consistent, trusted, well-timed input data arriving from the integrated systems.
- C. Isolated app testing is technically impossible for AI features.
- D. Dynamics 365 apps cannot exchange data without a manual export/import step.

Correct Answer: B

Why: Copilot and agent responses are grounded in data pulled from multiple systems. A defect in the data flow between apps (staleness, missing fields, timing) degrades AI output quality even when each app functions correctly on its own — hence end-to-end, cross-app scenario testing is required.

Q4. MULTIPLE CHOICE — SINGLE ANSWER

In the Copilot test-case blueprint template, which field is specifically intended to capture inputs such as an empty report, a corrupted PDF, or a missing revenue section?

- A. Preconditions

- B. Dependencies
- C. Edge Case Variations
- D. Expected Results

Correct Answer: C

Why: The blueprint's 'Edge Case Variations' field exists to document boundary conditions and abnormal inputs (empty/corrupted/missing data) so Copilot-generated tests explicitly cover failure paths, not just the happy path.

Q5. MULTIPLE CHOICE — MULTIPLE ANSWERS

When an architect writes a standardized prompt pattern to instruct Copilot to generate test cases, which elements should the prompt include? (Select all that apply.)

- A. Goal — what Copilot should produce
- B. Context — business rules, code, or workflow descriptions
- C. Constraints — format, tone, coverage depth, granularity
- D. Quality Expectations — definition of 'good' output
- E. The architect's personal calendar availability

Correct Answer: A, B, C, D

Why: Effective Copilot prompting for test-case generation combines Goal, Context, Constraints, and Quality Expectations so the model produces predictable, high-quality, on-scope test cases. Calendar availability is irrelevant to prompt design.

Q6. MULTIPLE CHOICE — SINGLE ANSWER

On the Copilot test maturity ladder, what distinguishes Level 5 from Level 3?

- A. Level 5 uses ad hoc prompting while Level 3 is fully governed.
- B. Level 5 establishes enterprise-wide prompt governance and reusable test frameworks, while Level 3 only automates prompt pipelines integrated with CI/CD.
- C. Level 5 and Level 3 are functionally identical.
- D. Level 5 removes the need for CI/CD integration entirely.

Correct Answer: B

Why: The maturity ladder progresses from ad hoc prompting (L1) through standard templates (L2), automated CI/CD pipelines (L3), continuous risk-based/dynamic regeneration testing (L4), to enterprise-wide governance with reusable frameworks (L5) — the highest maturity level.

Q7. DROP-DOWN (SELECT FROM LIST)

Complete the Copilot test-case blueprint by selecting the correct field for each description.

Blank	Drop-down options	Correct selection
Blank 1	Test ID / Dependencies / Steps	Test ID — unique identifier such as TC-001-FIN-Q1
Blank 2	Preconditions / Expected Results / Purpose/Description	Preconditions — required setup such as an uploaded
Blank 3	Dependencies / Inputs/Parameters / Edge Case Variations	Dependencies — external requirements such as kno

Why: The blueprint separates identity (Test ID), setup state (Preconditions), and external requirements (Dependencies) from execution data — mapping the correct field to each description keeps generated test cases consistent and machine-parsable.

Q8. MULTIPLE CHOICE — SINGLE ANSWER

During validation of Copilot-generated test cases, which of the following is NOT one of the four review dimensions defined in the strategy?

- A. Completeness
- B. Accuracy
- C. Cost efficiency
- D. Maintainability

Correct Answer: C

Why: The defined review dimensions are Completeness, Accuracy, Clarity, and Maintainability. Cost efficiency is a valid operational concern elsewhere in ALM but is not one of the four Copilot test-case validation criteria.

Q9. DRAG AND DROP

Arrange the stages of the Copilot-driven test-case creation strategy map in the correct order.

Draggable items: Architect Validation Pass | Business Requirements | Copilot Test Case Drafts | Final Test Case Library | Architect Prompt Plan

Targets (in order): Stage 1 → Stage 2 → Stage 3 → Stage 4 → Stage 5

Target	Correct item placed
Stage 1	Business Requirements
Stage 2	Architect Prompt Plan
Stage 3	Copilot Test Case Drafts
Stage 4	Architect Validation Pass
Stage 5	Final Test Case Library

Why: The pipeline flows from Business Requirements into an Architect Prompt Plan (Goal/Context/Constraints/Blueprint), which Copilot uses to produce Test Case Drafts; the architect then performs a Validation Pass (completeness/accuracy/edge cases/compliance) before publishing the Final Test Case Library.

Q10. MULTIPLE CHOICE — MULTIPLE ANSWERS

Which activities belong under the 'Completeness' review criterion when validating Copilot-generated test cases? (Select all that apply.)

- A. All acceptance criteria covered
- B. Key workflows represented
- C. Boundary and failure conditions included
- D. Steps are actionable and unambiguous
- E. Naming conventions are consistent

Correct Answer: A, B, C

Why: Completeness concerns coverage — acceptance criteria, workflows, and boundary/failure conditions. Actionable/unambiguous steps fall under Clarity, and naming conventions fall under Maintainability.

Q11. MULTIPLE CHOICE — SINGLE ANSWER

A finance application's business rules change mid-quarter. What should trigger Copilot to regenerate the affected test cases as part of a scalable testing strategy?

- A. A manually scheduled annual review only
- B. A regression trigger tied to the requirements change, using the versioned prompt library and templates
- C. Waiting until a production incident occurs

D. Disabling testing until the next major release

Correct Answer: B

Why: Scaling a Copilot-driven strategy requires a reusable, version-controlled prompt library plus regression triggers so that when requirements change, Copilot is proactively re-invoked to generate updated tests — rather than waiting for an incident or an arbitrary annual cycle.

Q12. HOTSPOT / CASE STUDY (YES / NO STATEMENTS)

Scenario: An architect reviews a Copilot-generated regression suite for a Q1 financial summarization agent before sign-off. The suite includes 40 test cases covering standard inputs, two edge cases (empty report, corrupted PDF), and reused Test IDs from a prior release without modification.

Statement	Yes / No
The reused, unmodified Test IDs from a prior release satisfy the Maintainability criterion.	No
Including the empty-report and corrupted-PDF cases satisfies part of the Completeness criterion.	Yes
The architect can approve the suite for production without checking whether expected results are measurable.	No

Why: Reusing Test IDs without validating they still map correctly to current logic risks naming/versioning drift and does not by itself satisfy Maintainability (which requires consistent, version-correct reuse). Edge cases do satisfy part of Completeness. Expected results must be measurable and reviewed for Clarity before sign-off — skipping that check is not acceptable.

Section B — Design ALM Process for AI-Powered Business Solutions (Q13–Q40)

Q13. MULTIPLE CHOICE — SINGLE ANSWER

Under an AI data ALM strategy, how should training/fine-tuning datasets, evaluation ('golden') sets, grounding knowledge, and prompt assets be treated?

- A. As disposable working files not subject to change control
- B. As versioned, promotable ALM artifacts with defined owners and gates
- C. As Production-only assets exempt from Dev/Test governance
- D. As items that only Security & Compliance may version

Correct Answer: B

Why: The data ALM model explicitly scopes datasets, grounding knowledge, prompt assets, policies/guardrails, and run telemetry as versioned, promotable artifacts — each requiring ownership and gate evidence before promotion, just like code.

Q14. MULTIPLE CHOICE — SINGLE ANSWER

What does the 'red/gold' dataset pattern describe in an AI data environment strategy?

- A. Red data is production-approved; gold data is experimental.
- B. Red data is mutable and experimental; gold data is frozen and promoted.
- C. Red and gold refer to sensitivity labels only, unrelated to promotion state.
- D. Red data is used only for UI testing; gold data is used only for load testing.

Correct Answer: B

Why: Red/gold is a promotion-state pattern: red datasets are mutable and used for active experimentation, while gold datasets are frozen, immutable, and have passed the promotion gates required for advancement toward production.

Q15. DROP-DOWN (SELECT FROM LIST)

Match each ALM promotion gate to the evidence it primarily requires.

Blank	Drop-down options	Correct selection
Gate A→B	Data contract; catalog/ownership policy proof / Metrics pack; safety card; catalog/ownership policy proof	Data contract; catalog/ownership policy proof
Gate C→D	Profiling report; lineage IDs / Metrics pack; safety card / Residency mapping pack; safety memo	Metrics pack; safety memo
Gate E→F	Data contract; catalog policy proof / Residency mapping; approval dashboard; dashboard; budget guard	Dashboard; dashboard; budget guard

Why: Gate A→B verifies ownership/labeling via the data contract; Gate C→D verifies evaluation and safety via a metrics pack and safety card; Gate E→F verifies runtime readiness via dashboards, alarms, and cost guardrails before entering full operation.

Q16. MULTIPLE CHOICE — MULTIPLE ANSWERS

Which activities occur specifically during Phase D — 'Stage & Approve' — of the AI data ALM lifecycle? (Select all that apply.)

- A. Execute canary runs using masked/representative Prod-like data
- B. Freeze gold datasets and sign immutability attestations
- C. Perform privacy, security, and compliance reviews (DLP, RAI, export controls)
- D. Register the release in the data catalog and publish consumer contracts

E. Train or iterate the model using Dev/Test data

Correct Answer: A, B, C

Why: Phase D covers pre-production staging and approval: compliance reviews, canary runs on masked Prod-like data, and freezing/attesting gold datasets. Registering the release and publishing consumer contracts happens in Phase E (Deploy & Serve); training/iterating happens in Phase C (Develop & Evaluate).

Q17. MULTIPLE CHOICE — SINGLE ANSWER

How is canary testing best described in the context of AI data ALM?

- A. A full production rollout to all users simultaneously
- B. A low-risk deployment strategy that releases new code to a small, isolated subset of users or servers to identify issues before full rollout
- C. A method for generating synthetic training data only
- D. A compliance audit performed after go-live

Correct Answer: B

Why: Canary testing limits exposure by releasing changes to a small, isolated subset first, allowing issues to surface with limited blast radius before a full production rollout.

Q18. DRAG AND DROP

Match each RACI activity to the role that is Accountable (A) per the representative data ALM RACI chart.

Draggable items: AI Architect | Data Owner | Security/Compliance | Platform Admin

Targets (in order): Define data contract → Evaluation thresholds → Residency configuration → Retention & deletion

Target	Correct item placed
Define data contract	Data Owner
Evaluation thresholds	AI Architect
Residency configuration	Security/Compliance
Retention & deletion	Data Owner

Why: Per the sample RACI: Data Owner is Accountable/Responsible for defining the data contract and for retention & deletion; AI Architect is Accountable (with Responsible) for evaluation thresholds; Security/Compliance is Accountable for residency configuration. Roles should always be validated per implementation.

Q19. MULTIPLE CHOICE — SINGLE ANSWER

For a regulated workload, what should the default cross-border data processing posture be for Copilot/Power Platform features?

- A. Enable cross-region overflow by default for performance
- B. Default to in-region processing and require explicit approval to enable overflow
- C. Randomly distribute workloads across regions to balance load
- D. Disable all AI features until a global data center is available

Correct Answer: B

Why: In regulated scenarios, the residency decision tree sets the default to in-region processing; cross-region overflow is only enabled under explicit admin-controlled approval, or the feature is blocked/deferred if overflow isn't permitted for that workload tier.

Q20. MULTIPLE CHOICE — SINGLE ANSWER

In the recommended Copilot Studio environment strategy, where are managed solutions required?

- A. Only in the Development environment
- B. In Test and Production, not in Development
- C. Managed solutions are never used for Copilot Studio
- D. Only when connectors are involved

Correct Answer: B

Why: Dev remains unmanaged for iterative authoring, while Test and Production use managed solutions only, enforcing no direct editing in production and consistent, controlled promotion.

Q21. MULTIPLE CHOICE — SINGLE ANSWER

In the Connector Release Flow (Author → Validate Auth → Apply DLP → Approve Security → Publish to Prod), what immediately follows 'Apply DLP'?

- A. Author
- B. Validate Auth
- C. Approve Security
- D. Publish to Prod

Correct Answer: C

Why: The defined release flow order is Author, then Validate Auth, then Apply DLP, then Approve Security, and finally Publish to Prod — DLP application must be confirmed before the security approval gate.

Q22. DRAG AND DROP

Arrange the connector ALM release steps in the correct sequence.

Draggable items: Approve Security | Author | Publish to Prod | Validate Auth | Apply DLP

Targets (in order): Step 1 → Step 2 → Step 3 → Step 4 → Step 5

Target	Correct item placed
Step 1	Author
Step 2	Validate Auth
Step 3	Apply DLP
Step 4	Approve Security
Step 5	Publish to Prod

Why: Connectors must be authored in Dev, have authentication flows validated, have DLP policies applied during Test, pass security approval, and only then be published to Production — skipping or reordering steps breaks the governance chain.

Q23. MULTIPLE CHOICE — SINGLE ANSWER

What versioning principle should be applied to Microsoft Foundry agents under ALM?

- A. Agents should be treated as disposable and never versioned.
- B. Each agent should be treated as a versioned artifact with immutable production releases and documented changelogs.
- C. Only the underlying model needs versioning, not the agent configuration.
- D. Version history should be discarded after each production promotion.

Correct Answer: B

Why: Foundry agent ALM requires treating each agent as a versioned artifact, maintaining immutable releases for production, documenting changes via changelogs, and promoting only versioned bundles across environments.

Q24. MULTIPLE CHOICE — SINGLE ANSWER

Which of the following is checked at Foundry Promotion Gate 1 (Dev → Test)?

- A. Regression test completion and cost assessment
- B. Human validation of agent reasoning
- C. Functional validation of core behaviors and initial safety/guardrail checks
- D. Production release calendar sign-off

Correct Answer: C

Why: Gate 1 (Dev→Test) focuses on early validation: functional behavior checks, initial safety/guardrail checks, data source mapping verification, and action/connector configuration review. Regression completion and human reasoning validation belong to Gate 2 (Test→Prod).

Q25. MULTIPLE CHOICE — SINGLE ANSWER

At Foundry Promotion Gate 2 (Test → Prod), which requirement specifically distinguishes it from Gate 1?

- A. Initial safety and guardrail checks
- B. Human validation of agent reasoning, plus documented version/dependency/risk analysis
- C. Code or prompt quality review
- D. Verification of data source mappings

Correct Answer: B

Why: Gate 2 adds regression completion, data-policy compliance approval, performance/cost assessment, human validation of reasoning, and documentation of version/dependencies/risk — a deeper bar than the initial functional/safety checks performed at Gate 1.

Q26. MULTIPLE CHOICE — MULTIPLE ANSWERS

Which components must be governed as ALM-managed artifacts for a Microsoft Foundry agent? (Select all that apply.)

- A. Agent logic and orchestration
- B. Prompts and behavioral rules
- C. Action handlers and workflows
- D. Security and permission mappings
- E. The end user's personal device settings

Correct Answer: A, B, C, D

Why: Foundry agent components requiring ALM governance include logic/orchestration, prompts/behavioral rules, action handlers/workflows, data/grounding configuration, security/permission mappings, connectors, and policy/safety controls. End-user device settings are out of scope.

Q27. MULTIPLE CHOICE — SINGLE ANSWER

Which artifact captures a custom AI model's metrics, constraints, and recommended usage boundaries before it can be approved for deployment?

- A. RACI matrix
- B. Model Card
- C. Data Contract
- D. Connector manifest

Correct Answer: B

Why: The Evaluation & Approval stage of custom-model ALM produces a Model Card documenting accuracy/reliability metrics, cost/latency, safety findings, and recommended/constrained usage boundaries.

Q28. DROP-DOWN (SELECT FROM LIST)

Select the correct promotion-gate criteria for custom AI models.

Blank	Drop-down options	Correct selection
Gate 1 (Dev→Test)	Training data meets quality standards; safety/bias testing passed	Training data meets quality standards; safety/bias testing passed
Gate 2 (Test→Prod)	Training data meets quality standards only / Documentation created	Performance exceeds baseline; cost/latency acceptable / Rollback plan confirmed

Why: Gate 1 confirms the model is fundamentally sound (data quality, safety/bias testing, documentation) before entering Test. Gate 2 confirms it is production-ready — performance, cost, security/compliance sign-off, and a confirmed rollback plan — before promotion to Production.

Q29. MULTIPLE CHOICE — SINGLE ANSWER

What does 'BYOM' refer to in the custom AI model development stage?

- A. Build Your Own Metrics
- B. Bring Your Own Model — incorporating externally built or fine-tuned model components
- C. Backup Your Operational Memory
- D. Batch Your Output Messages

Correct Answer: B

Why: BYOM ('bring your own model') refers to incorporating externally developed or pre-existing model components into a solution rather than training entirely from scratch, while still applying responsible AI and ALM governance to it.

Q30. HOTSPOT / CASE STUDY (YES / NO STATEMENTS)

Scenario: A team wants to promote a fraud-detection custom AI model from Test to Production. Its accuracy exceeds baseline, latency is within budget, and a rollback plan exists, but the security and compliance team has not yet signed off on the deployment package.

Statement	Yes / No
The model satisfies all Gate 2 (Test→Prod) criteria and can be promoted immediately.	No
The Model Card should document the model's constraints and recommended usage boundaries prior to approval.	Yes
Rollback plan confirmation alone is sufficient evidence for the Gate 2 promotion decision.	No

Why: Gate 2 requires performance, cost/latency, security & compliance approval, AND a rollback plan — missing security/compliance sign-off means the gate is not satisfied even though other criteria pass. The Model Card is a required artifact for the Evaluation & Approval stage.

Q31. MULTIPLE CHOICE — SINGLE ANSWER

In the Dynamics 365 Finance and Supply Chain (F&SC;) AI environment strategy, what is the primary purpose of the TEST environment?

- A. Execute approved AI capabilities in live financial workloads
- B. Validate with safe, anonymized production-like data and perform regression checks
- C. Build and iterate AI models, prompts, and orchestration logic from scratch
- D. Store only synthetic marketing content

Correct Answer: B

Why: TEST is where anonymized, production-like data is used to validate models and perform regression checks, bridging DEV (build/iterate) and PROD (execute live workloads).

Q32. MULTIPLE CHOICE — SINGLE ANSWER

Which key ALM requirement applies uniformly across DEV/TEST/PROD for AI in Dynamics 365 F&SC;?

- A. Direct edits are allowed in Production if approved verbally.
- B. No direct edits in production; environment variables are used for endpoints, model names, and connection strings.
- C. Only PROD requires environment variables; DEV and TEST can hardcode values.
- D. Synthetic datasets are prohibited in TEST.

Correct Answer: B

Why: Key requirements include prohibiting direct production edits and using environment variables (rather than hardcoded values) for endpoints, model names, and connection strings — enabling safe, repeatable promotion across environments.

Q33. MULTIPLE CHOICE — SINGLE ANSWER

In the Dynamics 365 Customer Service/Customer Engagement AI ALM workflow, which environment is used to validate business acceptance, performance, safety, and compliance with live-like customer service queues before Production?

- A. Development (DEV)
- B. Test/Validation (TEST)
- C. Pre-Production (UAT/PREPROD)
- D. Production (PROD)

Correct Answer: C

Why: Pre-Production (UAT/PREPROD) is the stage dedicated to business acceptance, performance, safety, and compliance validation with live-like interactions and knowledge entities, ahead of full Production rollout.

Q34. MULTIPLE CHOICE — MULTIPLE ANSWERS

Which items should be included as AI assets under ALM governance for Dynamics 365 Customer Service/CE AI capabilities? (Select all that apply.)

- A. Copilot prompts and agent behaviors
- B. Knowledge sources and semantic indexing configurations
- C. Power Automate flows and plugins used by AI actions
- D. Environment variable sets (API keys, model endpoints, vector index references)
- E. The CRM sales team's personal email signatures

Correct Answer: A, B, C, D

Why: ALM for D365 CE AI governs prompts/behaviors, knowledge sources/semantic indexing, Power Automate flows/plugins tied to AI actions, data contracts, and environment variable sets. Personal email signatures are unrelated business content, not an ALM asset.

Q35. MULTIPLE CHOICE — SINGLE ANSWER

Which release-versioning approach is recommended when marking AI asset releases for Dynamics 365 CE (e.g., v1.3.2)?

- A. Sequential integers only (1, 2, 3)
- B. Semantic versioning
- C. Date-stamp only (YYYYMMDD)

D. Random build hashes

Correct Answer: B

Why: The workflow explicitly recommends marking releases with semantic versioning (major.minor.patch, e.g., v1.3.2) to communicate the scope of change consistently across teams.

Q36. MULTIPLE CHOICE — SINGLE ANSWER

Which dataset is most ready for promotion to Production in an AI ALM process?

- A. The latest Dev extract with broad sample coverage
- B. A Test dataset that looks balanced but has partial lineage
- C. An immutable, versioned dataset with sensitivity labels and full lineage that passed evaluation gates
- D. Any dataset that delivers acceptable latency in a pilot

Correct Answer: C

Why: Production promotion requires the dataset to be immutable, versioned, properly labeled for sensitivity, fully lineage-traced, and to have passed the evaluation/quality gates — partial lineage, unreviewed Dev extracts, or latency-only evidence are all insufficient.

Q37. DROP-DOWN (SELECT FROM LIST)

Match each Dynamics 365 F&SC; environment to its correct purpose.

Blank	Drop-down options	Correct selection
DEV	Build and iterate AI models, prompts, orchestration logic, and integrations	Build and Iterate
PROD	Build and iterate AI models / Validate with anonymized production-like data	Execute Approved

Why: DEV is for building/iterating; TEST validates with safe anonymized production-like data; PROD executes approved capabilities against live workloads — this separation prevents untested logic from reaching regulated financial processes.

Q38. MULTIPLE CHOICE — SINGLE ANSWER

Which Gate in the Dynamics 365 F&SC; AI ALM process requires the deployment package to be approved by both architecture and governance teams?

- A. Gate 1 — DEV → TEST
- B. Gate 2 — TEST → PROD
- C. There is no such requirement
- D. It occurs only during retirement

Correct Answer: B

Why: Gate 2 (TEST→PROD) criteria include model evaluation thresholds met, no bias/unsafe output, verified ERP workload performance, completed security/compliance/residency reviews, and explicit deployment package approval by architecture and governance teams.

Q39. MULTIPLE CHOICE — MULTIPLE ANSWERS

Which items belong on the AI Release Readiness Checklist for Dynamics 365 Customer Service/CE AI features? (Select all that apply.)

- A. Data quality validated
- B. Knowledge sources aligned and tested
- C. Prompts regression tested
- D. Monitoring dashboards ready

E. Marketing budget finalized

Correct Answer: A, B, C, D

Why: The checklist covers data quality, knowledge source alignment/testing, prompt regression testing, safety/compliance verification, environment variable configuration, deployment pipeline success, and monitoring readiness. Marketing budget is unrelated to AI release readiness.

Q40. MULTIPLE CHOICE — SINGLE ANSWER

Under the red/gold dataset pattern, which dataset state must be reached before a dataset can be used to power a Production AI feature?

- A. Red — mutable, still being experimented on
- B. Gold — frozen and formally promoted
- C. Either state is acceptable for Production
- D. Neither — Production always trains live on incoming data

Correct Answer: B

Why: Only gold (frozen, promoted, immutable, evidence-backed) datasets should power Production AI features; red datasets remain in active experimentation and are not yet promotion-ready.

Section C — Design Responsible AI Security, Governance, Risk Management & Compliance (Q41–Q80)

Q41. MULTIPLE CHOICE — SINGLE ANSWER

What is the primary purpose of assigning a unique managed identity to each AI agent per environment (prod, pre-prod, dev)?

- A. To simplify billing reports only
- B. To remove embedded secrets and enable least-privilege, environment-scoped authentication with clear ownership
- C. To allow the agent to impersonate any user without restriction
- D. To bypass RBAC checks for faster execution

Correct Answer: B

Why: Managed identities eliminate embedded secrets, simplify credential rotation, and let each agent instance (per environment) be scoped to least-privilege roles with recorded ownership and lifecycle metadata.

Q42. MULTIPLE CHOICE — SINGLE ANSWER

When an agent acts on behalf of a signed-in user versus acting as itself, what authorization pattern should be applied?

- A. Always use the agent's own broad service role regardless of context
- B. Propagate the user's permissions when acting on their behalf; use a narrowly scoped service role limited to needed actions when acting as itself
- C. Grant the agent global admin rights to simplify testing
- D. Disable authorization checks for internal agents

Correct Answer: B

Why: Correct authorization design distinguishes delegated action (propagate the user's own permissions) from autonomous action (use a minimally scoped service role) — both governed by least privilege.

Q43. MULTIPLE CHOICE — MULTIPLE ANSWERS

Which roles should be kept distinct under separation-of-duties for agent governance? (Select all that apply.)

- A. Maker
- B. Publisher
- C. Environment Admin
- D. Security Admin
- E. End customer

Correct Answer: A, B, C, D

Why: Separation of duties requires distinct roles for Maker, Publisher, Environment Admin, and Security Admin, with approvals required for production publishing and high-risk capability changes. The end customer is a consumer of the agent, not a governance role.

Q44. MULTIPLE CHOICE — SINGLE ANSWER

What is the recommended approach to keep confidential data sources out of a public-facing customer agent while satisfying data residency requirements?

- A. Merge all internal and public data into one workload for simplicity
- B. Separate internal vs. public workloads and select compliant regions for knowledge, logs, and memory

- C. Store all agent memory in a single global region regardless of source sensitivity
- D. Skip data boundary design since agents can self-regulate access

Correct Answer: B

Why: Data governance for agents requires separating internal and public-facing workloads (keeping confidential sources out of public agents) and honoring data residency by choosing compliant regions for knowledge, logs, and memory storage.

Q45. MULTIPLE CHOICE — SINGLE ANSWER

Which control specifically restricts connectors, actions, and data movement to prevent sensitive data from leaving approved boundaries?

- A. Sensitivity labels alone
- B. DLP (data loss prevention) policies
- C. Managed identity rotation
- D. Red-team exercises

Correct Answer: B

Why: DLP policies are the mechanism used to restrict which connectors and actions an agent may use and to prevent unauthorized data movement, complementing (not replacing) sensitivity labeling.

Q46. MULTIPLE CHOICE — MULTIPLE ANSWERS

Which activities fall under 'Observability and cost governance' for AI agents? (Select all that apply.)

- A. Centralizing telemetry for prompts, tool calls, errors, and safety events
- B. Maintaining an authoritative agent inventory with owner, version, environment, and purpose
- C. Tagging resources by agent/cost center and setting spend alerts
- D. Writing the agent's marketing copy
- E. Capturing custom business metrics such as task completion and escalation rates

Correct Answer: A, B, C, E

Why: Observability/cost governance includes unified logging, an agent inventory/catalog, cost tagging with spend alerts, and business metric capture. Marketing copy is unrelated to observability.

Q47. MULTIPLE CHOICE — SINGLE ANSWER

What is the goal of running adversarial ('red team') evaluations for AI agents?

- A. To measure UI responsiveness under load
- B. To reduce the blast radius of adversarial input and model-specific risks by testing pre-production and after major changes
- C. To replace unit testing entirely
- D. To validate billing accuracy

Correct Answer: B

Why: Red team evaluations proactively probe for adversarial and model-specific vulnerabilities before release and after major changes, with release gating on the closure of identified findings.

Q48. MULTIPLE CHOICE — SINGLE ANSWER

Which protocol is recommended for structured tool/data access by AI agents, as distinct from controlled delegation between agents?

- A. Model Context Protocol (MCP) for tool/data access; Agent-to-Agent (A2A) for delegation
- B. HTTP/REST is sufficient for both use cases; no standard protocol is needed

- C. Agent-to-Agent (A2A) for tool/data access; MCP for delegation
- D. Both use the same protocol interchangeably with no distinction

Correct Answer: A

Why: MCP standardizes structured tool/data access for an agent, while A2A governs controlled delegation and context sharing across multiple agents — using the wrong protocol for the wrong purpose undermines the intended governance boundary.

Q49. DROP-DOWN (SELECT FROM LIST)

Select the correct RBAC scope descriptor for each governance role in the design checklist.

Blank	Drop-down options	Correct selection
Agent Runtime	Least privilege / Full admin access / Publisher-only rights	Least privilege
Publisher	Least privilege on Azure Roles / Unrestricted Azure subscription owner permissions / Azure Roles	Least privilege on Azure Roles

Why: The RBAC design table applies least-privilege scoping across both Agent Runtime and Publisher roles at the Azure platform layer, avoiding broad or unrestricted access even for roles that need to promote changes.

Q50. MULTIPLE CHOICE — SINGLE ANSWER

During incident response planning for AI agents, which capability must be defined in advance for a critical, compromised agent?

- A. How to disable the agent quickly, preserve logs, notify stakeholders, and recover safely
- B. How to permanently delete all logs immediately to save storage costs
- C. How to silently continue operation while investigating
- D. How to escalate only after a public disclosure occurs

Correct Answer: A

Why: Incident response for agents requires a predefined rapid-disable capability, log preservation for forensics, stakeholder notification, and a safe recovery path — logs must be preserved, not deleted, and operation should not silently continue during a suspected compromise.

Q51. MULTIPLE CHOICE — SINGLE ANSWER

Which statement best describes 'accountability and ownership' as a governance principle for AI agents?

- A. Ownership is optional if the agent operates only internally.
- B. An agent owner is assigned responsibility for lifecycle, security posture, and approvals, and agents are documented in a registry.
- C. Accountability transfers automatically to Microsoft once an agent is deployed on Azure.
- D. Only the security team needs to know an agent exists.

Correct Answer: B

Why: Accountability requires a named agent owner responsible for lifecycle/security/approvals and a registry documenting purpose, environment, risk level, and data access for every agent — organizational responsibility does not transfer to the cloud provider.

Q52. MULTIPLE CHOICE — MULTIPLE ANSWERS

Which layers appear in the Data Governance Layers model for agents? (Select all that apply.)

- A. Data Sources
- B. Agent Runtime
- C. DLP & Sensitivity Filters

- D. User Output & Logging
- E. Sales Forecast Dashboard

Correct Answer: A, B, C, D

Why: The Data Governance Layers flow is Data Sources → Agent Runtime → DLP & Sensitivity Filters → User Output & Logging, ensuring data is filtered for sensitivity and DLP before it becomes agent output, and that output is logged for audit. A sales dashboard is not part of this governance layering.

Q53. MULTIPLE CHOICE — SINGLE ANSWER

Which technique is a recommended safeguard specifically against model drift and data poisoning during retraining pipelines?

- A. Skipping data lineage tracking to speed up retraining
- B. Tracking data lineage for all training sources and validating new data before use in retraining
- C. Allowing any external dataset to be merged into training data without review
- D. Disabling anomaly detection to reduce noise in logs

Correct Answer: B

Why: Guarding against drift/poisoning requires lineage tracking for training sources, validation of newly introduced data before retraining, and automated anomaly detection on model behavior/output patterns — the opposite of unreviewed data merges.

Q54. MULTIPLE CHOICE — SINGLE ANSWER

In the model hardening blueprint, what is the correct order of the five components?

- A. Monitoring → Validation → Threat Protection → Private Endpoints → Secure Compute
- B. Secure Compute → Private Endpoints → Threat Protection → Validation Pipeline → Monitoring and Drift Detection
- C. Threat Protection → Secure Compute → Monitoring → Validation → Private Endpoints
- D. Private Endpoints → Monitoring → Secure Compute → Threat Protection → Validation

Correct Answer: B

Why: The blueprint proceeds from foundational infrastructure controls (Secure Compute, Private Endpoints) through Threat Protection, then a Validation Pipeline for models/data, and finally ongoing Monitoring and Drift Detection.

Q55. MULTIPLE CHOICE — SINGLE ANSWER

Which practice best reflects data minimization for a model handling customer records?

- A. Ingest all available fields in case they're useful later
- B. Use only the minimum data required for the model's purpose and redact PII/secrets/financial identifiers during preprocessing
- C. Store raw customer records permanently for auditability
- D. Allow the model to request any field directly from source systems at inference time

Correct Answer: B

Why: Data minimization restricts a model to the minimum data necessary for its stated purpose and requires redaction of sensitive fields during preprocessing — broad, unrestricted ingestion increases exposure risk without a corresponding business need.

Q56. MULTIPLE CHOICE — MULTIPLE ANSWERS

Which are recognized categories of prompt manipulation risk? (Select all that apply.)

- A. Overriding system instructions (e.g., 'ignore previous instructions...')
- B. Deceptive context claiming false authorization to disclose confidential data

- C. Embedding hidden instructions in text, HTML, or files
- D. Multi-step coercion or poisoning inputs
- E. Increasing the model's context window size

Correct Answer: A, B, C, D

Why: Prompt manipulation techniques include instruction overriding, deceptive/false-authorization context, hidden embedded instructions, and multi-step coercion/poisoning. Context window size is a technical model parameter, not a manipulation technique.

Q57. MULTIPLE CHOICE — SINGLE ANSWER

Which mitigation category is most directly aimed at reducing identity/RBAC exploitability in AI systems?

- A. Strengthening input and output filtering only
- B. Enforcing least-privilege access, using managed identities, segmenting dev/test/prod, and reviewing access logs for anomalous elevation
- C. Increasing model temperature to diversify outputs
- D. Disabling all logging to reduce attack surface

Correct Answer: B

Why: Reducing identity/RBAC exploitability specifically requires least-privilege enforcement, managed identities, environment segmentation, and access log review for anomalous privilege elevation — filtering addresses a different vulnerability category (prompt/output risk).

Q58. MULTIPLE CHOICE — SINGLE ANSWER

An agent workflow autonomously calls a non-Microsoft external endpoint without an approval step. Which vulnerability category does this best represent?

- A. Data exposure vulnerability only
- B. Agent and workflow-level vulnerability — autonomous tool use without proper guardrails
- C. Fairness and bias vulnerability
- D. Transparency vulnerability

Correct Answer: B

Why: Autonomous tool use without guardrails, misinterpreted instructions triggering unintended actions, poor auditing, and unsecured calls to non-Microsoft endpoints are all agent/workflow-level vulnerabilities requiring clearly bounded tool capabilities and end-to-end auditing.

Q59. DROP-DOWN (SELECT FROM LIST)

Select the correct mitigation for each AI vulnerability category.

Blank	Drop-down options	Correct selection
Data exposure vulnerability	Minimization, RBAC boundaries, intent-based access / Increased model size / Disabling intent-based access	Data minimization / RBAC / Intent-based access
Model behavior vulnerability	Continuous evaluation output safety, correctness, consistency, and drift / Only evaluating post-release / Removing the model	Continuous evaluation of outputs / Removing the model

Why: Data exposure risks are mitigated through minimization and access boundaries aligned to user roles; model behavior risks (fabrication, toxicity, bias, drift) are mitigated through continuous structured evaluation of outputs across releases.

Q60. MULTIPLE CHOICE — SINGLE ANSWER

Which six principles make up Microsoft's Responsible AI framework referenced in the RAI solution review?

- A. Fairness, Reliability and Safety, Privacy and Security, Inclusiveness, Transparency, Accountability
- B. Speed, Cost, Scalability, Usability, Aesthetics, Popularity

- C. Confidentiality, Integrity, Availability, Non-repudiation, Authentication, Authorization
- D. Accuracy, Precision, Recall, F1 Score, Latency, Throughput

Correct Answer: A

Why: The six Responsible AI principles are Fairness, Reliability and Safety, Privacy and Security, Inclusiveness, Transparency, and Accountability — these form the lens for evaluating models, agents, workflows, and user experiences.

Q61. MULTIPLE CHOICE — SINGLE ANSWER

When reviewing 'Model and agent behavior evaluation' as part of a Responsible AI review, which task is explicitly included?

- A. Confirming fallback behavior for unknown or ambiguous requests
- B. Approving the marketing budget for the launch
- C. Selecting the UI color palette
- D. Negotiating vendor contracts for compute

Correct Answer: A

Why: Model/agent behavior evaluation includes identifying fabricated information, validating instruction/prompt adherence, assessing safety boundaries in edge cases, confirming fallback behavior for unknown/ambiguous requests, and ensuring outputs are not harmful or misleading.

Q62. MULTIPLE CHOICE — MULTIPLE ANSWERS

Which stages appear in the Responsible AI Lifecycle diagram? (Select all that apply, in the order they should be read left to right conceptually.)

- A. Design
- B. Build
- C. Validate
- D. Deploy
- E. Ignore feedback

Correct Answer: A, B, C, D

Why: The Responsible AI Lifecycle progresses through Design → Build → Validate → Deploy (and continues into ongoing monitor/operate stages). Ignoring feedback contradicts the RAI principle of continuous monitoring.

Q63. MULTIPLE CHOICE — SINGLE ANSWER

Which governance mechanism enforces 'sunset criteria' for models that no longer meet safety or compliance requirements?

- A. A one-time design review with no follow-up
- B. Operational oversight and governance, including governance board reviews and continuous monitoring of logs, safety events, and drift
- C. A purely automated process with no human review
- D. The original developer's personal discretion

Correct Answer: B

Why: Sunset criteria are enforced through ongoing operational oversight — governance board reviews for significant updates/retraining and continuous evaluation of logs, safety events, and drift — not a single point-in-time review.

Q64. MULTIPLE CHOICE — SINGLE ANSWER

Which factor most directly determines where Copilot Studio prompt data and agent interactions are processed?

- A. The end user's device operating system
- B. The environment's regional configuration and whether unpublished agents/preview features follow different residency rules
- C. The color theme selected in the agent's UI
- D. The number of connectors configured

Correct Answer: B

Why: Copilot Studio enforces regional boundaries based on environment configuration; architects must specifically validate whether preview features or unpublished agents follow different (potentially less restrictive) residency rules than published production agents.

Q65. MULTIPLE CHOICE — SINGLE ANSWER

Which Microsoft Purview capability specifically prevents sensitive data from being used in AI prompts or outputs?

- A. Sensitivity labeling alone, without DLP
- B. Data loss prevention (DLP) rules
- C. Tenant region selection
- D. Backup and recovery configuration

Correct Answer: B

Why: Purview DLP rules are the control specifically applied to prevent sensitive data from being used in AI prompts or returned in outputs, complementing sensitivity labels which restrict cross-tenant/cross-region transmission.

Q66. MULTIPLE CHOICE — MULTIPLE ANSWERS

Which practices support a compliant AI architecture for data residency? (Select all that apply.)

- A. Selecting tenant regions aligned with regulatory frameworks
- B. Configuring Copilot Studio environments to enforce residency policies
- C. Allowing custom connectors to silently bypass regional data boundaries for convenience
- D. Documenting all data flows including logs, telemetry, and inference outputs
- E. Validating backup, recovery, and logging systems maintain compliance

Correct Answer: A, B, D, E

Why: A compliant architecture aligns tenant regions, enforces environment residency policy, documents all data flows, and validates backup/recovery/logging compliance. Allowing connectors to bypass regional boundaries directly violates data residency governance.

Q67. MULTIPLE CHOICE — SINGLE ANSWER

What design principle should govern access to grounding data used in retrieval-augmented generation (RAG)?

- A. Grant broad access to all grounding sources so retrieval never fails
- B. Apply least privilege by default, with role-aligned data partitions and clear ownership/stewardship
- C. Allow any authenticated user to query any grounding index without filtering
- D. Skip auditability requirements for read-only retrieval operations

Correct Answer: B

Why: Grounding data access must follow least privilege by default, partition data by function/role, assign accountable stewards per domain, and log all access events for auditability — broad or unfiltered access undermines confidentiality and governance.

Q68. MULTIPLE CHOICE — SINGLE ANSWER

Which step in the Retrieval Access Flow (Prompt → Policy Check → Search Index → Sanitization Layer → Model Context Injection) is responsible for enforcing DLP and sensitivity constraints before content reaches the model?

- A. Prompt
- B. Policy Check
- C. Sanitization Layer
- D. Model Context Injection

Correct Answer: C

Why: After the Policy Check authorizes the query and the Search Index returns candidates, the Sanitization Layer strips or filters disallowed content (per DLP/sensitivity rules) before the results are injected into the model's context.

Q69. MULTIPLE CHOICE — SINGLE ANSWER

Which requirement is specific to model-tuning access controls, distinguishing them from grounding-data retrieval access controls?

- A. Both require identical, unrestricted access for speed of iteration.
- B. Model tuning requires segregated environments, approval workflows for new training datasets, verified data lineage, and restricted ability to promote tuned models into production.
- C. Model tuning requires no approval workflow since it happens only in Dev.
- D. Retrieval access controls apply exclusively to tuning, not to RAG.

Correct Answer: B

Why: Model tuning (evaluation, fine-tuning, RL) requires elevated but tightly segregated permissions: dedicated environments, dataset-introduction approvals, verified lineage, security scanning of tuning inputs, and restricted promotion rights — insulating it from production.

Q70. MULTIPLE CHOICE — MULTIPLE ANSWERS

Which are examples of guardrails for grounding data and tuning inputs? (Select all that apply.)

- A. Block lists for prohibited document types
- B. Sanitization pipelines removing personal or contractual data
- C. Automated reviews validating safety and policy alignment
- D. Alerting and anomaly detection for unusual data access or tuning patterns
- E. Publishing all tuning datasets openly for transparency

Correct Answer: A, B, C, D

Why: Guardrails include block lists, sanitization pipelines, automated safety/policy reviews, and anomaly alerting. Publishing tuning datasets openly would violate confidentiality and access-control requirements, not serve as a guardrail.

Q71. MULTIPLE CHOICE — SINGLE ANSWER

Why should audit logs for AI model and data changes capture metadata rather than content?

- A. Metadata is cheaper to store but has no security rationale
- B. Capturing content would risk unnecessary exposure of sensitive information within the audit trail itself
- C. Content-based logs are technically impossible to generate
- D. Regulators require metadata-only logs in all jurisdictions without exception

Correct Answer: B

Why: Auditing data changes (ingestion, schema, refresh, labeling, movement) should capture metadata about the event rather than the sensitive content itself, avoiding the audit trail becoming a new source of data exposure.

Q72. MULTIPLE CHOICE — SINGLE ANSWER

Which architectural attribute is explicitly recommended for audit logs of model lifecycle events (registration, tuning, promotion, rollback, deployment)?

- A. Logs should be mutable so they can be corrected after the fact.
- B. Logs should be immutable, timestamped, role-attributed, and JSON-structured for machine parsing.
- C. Logs should be stored only in memory and discarded after each session.
- D. Logs should omit timestamps to reduce storage size.

Correct Answer: B

Why: Recommended audit attributes include immutability, timestamped change records, role-based attribution linked to identity, JSON-structured logs for machine parsing, and separation-of-duties/approval logging.

Q73. MULTIPLE CHOICE — SINGLE ANSWER

Which Azure AI Foundry capability specifically provides end-to-end visibility of model inference, identifying bottlenecks and correlating requests to model versions?

- A. Foundry activity logs only
- B. Foundry diagnostics and tracing
- C. RBAC role assignment blades
- D. Cost Management budgets

Correct Answer: B

Why: Foundry diagnostics and tracing provide traceability across model calls, pipeline activities, and tool interactions — including latency spikes, failure detection, and correlation of requests to specific model versions — beyond what activity logs alone capture.

Q74. DROP-DOWN (SELECT FROM LIST)

Select the appropriate retention pattern for each workload risk tier, per common audit-log retention guidance.

Blank	Drop-down options	Correct selection
Low-risk workload	90 days / 12–24 months / Indefinite retention	90 days
Regulated workload	90 days / 12–24 months / No retention required	12–24 months
Incident-related	90 days / 12–24 months / Indefinite retention	Indefinite retention

Why: Retention should be tiered by risk and defined jointly with Legal, Compliance, and InfoSec: shorter windows for low-risk workloads, extended windows for regulated workloads, and indefinite retention for incident-related archives supporting investigations.

Q75. DRAG AND DROP

Match each recommended tracing field to its purpose when designing audit pipelines with tracing.

Draggable items: Correlation ID | Model version | Response latency | Safety evaluation outcomes

Targets (in order): Links related events across a single request lifecycle → Attributes behavior/output differences to a specific release → Identifies performance bottlenecks → Confirms whether guardrails were satisfied for the response

Target	Correct item placed
Links related events across a single request lifecycle	Correlation ID
Attributes behavior/output differences to a specific release	Model version

Identifies performance bottlenecks	Response latency
Confirms whether guardrails were satisfied for the response	Safety evaluation outcomes

Why: Correlation ID ties together the events of one execution; model version attributes outcomes to a specific release; response latency exposes performance issues; safety evaluation outcomes confirm guardrail compliance for that specific response.

Q76. MULTIPLE CHOICE — SINGLE ANSWER

An architect is designing security for a helpdesk triage agent that reads ticket data, summarizes trends, and updates knowledge articles. Which identity design is correct?

- A. Use a single shared identity across Dev, Test, and Prod for simplicity
- B. Create per-environment managed identities with recorded ownership, and scope least-privilege Azure roles at the narrowest resource level
- C. Use the security administrator's personal account for all agent authentication
- D. Skip identity assignment since the agent only performs read operations

Correct Answer: B

Why: Even a read/summarize/update agent needs a dedicated, per-environment managed identity with documented ownership and narrowly scoped least-privilege roles — sharing identities across environments or using personal admin accounts both break the governance model.

Q77. HOTSPOT / CASE STUDY (YES / NO STATEMENTS)

Scenario: An enterprise deploys a customer-facing generative AI chatbot in Copilot Studio. The chatbot is configured with a public-facing knowledge source that also includes an internal confidential HR policy document. Logging is enabled centrally, and DLP policies restrict outbound connectors, but no sensitivity labels have been applied to the HR document.

Statement	Yes / No
Including the confidential HR document in a public-facing agent's knowledge source violates the recommended data boundary practice.	Yes
Enabled DLP policies alone are sufficient to fully compensate for the missing sensitivity labels on the HR document.	No
Centralized logging alone satisfies the full Responsible AI transparency requirement for this solution.	No

Why: Confidential sources must be kept out of public-facing agents entirely — DLP restricts connector/data movement but does not substitute for proper sensitivity labeling, and centralized logging is only one part (observability) of transparency, which also requires disclosing AI involvement and communicating limitations to users.

Q78. MULTIPLE CHOICE — SINGLE ANSWER

Which best reflects a defense-in-depth approach when designing enterprise AI agents?

- A. Grant broad access so the agent can retrieve any data it may need in future tasks
- B. Use layered identity, access, data governance, monitoring, and threat protection controls
- C. Allow the agent to self-correct risky behaviors without human oversight
- D. Disable logging to reduce operational costs

Correct Answer: B

Why: Defense-in-depth combines multiple independent layers — identity/access boundaries, data governance, observability, and threat protection — so that failure of any single control does not compromise the whole system. Broad access, unmonitored self-correction, and disabled logging each remove a necessary layer.

Q79. MULTIPLE CHOICE — SINGLE ANSWER

What is the most effective way to reduce the risk of AI agents exposing sensitive information?

- A. Allow unrestricted connector access to improve retrieval accuracy
- B. Rely solely on model instructions to avoid returning sensitive content
- C. Apply DLP policies, sensitivity labels, and least-privilege boundaries across all data sources
- D. Store sensitive data directly in agent prompts so it can reason more accurately

Correct Answer: C

Why: Prompt-level instructions alone are not a reliable security control since models can be manipulated; durable protection comes from layered technical controls — DLP, sensitivity labeling, and least-privilege data boundaries — not from broadening access or embedding sensitive data in prompts.

Q80. MULTIPLE CHOICE — MULTIPLE ANSWERS

Which practices together form a layered defense strategy for reducing AI-specific vulnerability risk, per the mitigation framework? (Select all that apply.)

- A. Identity and access boundaries
- B. Secure data governance
- C. Model evaluation and testing
- D. Monitoring and anomaly detection
- E. Publishing system prompts publicly so users can verify behavior

Correct Answer: A, B, C, D

Why: The layered-defense framework combines identity/access boundaries, secure data governance, structured model evaluation/testing, monitoring/anomaly detection, and incident response readiness. Publishing system prompts publicly is not a recommended control and can itself increase prompt-manipulation risk.