AB-100: Architect AI Solutions for Business Productivity

80 Advanced Practice Questions, Set 2 — Plan AI-Powered Business Solutions

Covers: ROAI & the Business Value Toolkit · data grounding & quality · prebuilt vs. custom Copilot Studio agents · multi-agent orchestration · AI-adoption governance & executive responsibilities · Responsible AI metrics · MCP server integration.

Structure: 10 single-select MCQ · 10 multi-select MCQ · 20 drag-and-drop · 20 drop-down · 2 case studies (10 sub-questions) · 3 hotspot Yes/No scenarios (10 statements). Difficulty: Architect / Advanced.

Part 1 — Single-Select Multiple Choice (Q1–Q10)

Q1. SINGLE-SELECT MULTIPLE CHOICE

An architect must calculate the return on AI investment (ROAI) for a proposed Copilot Studio agent by using the solution's usage metadata and telemetry. Which tool should the architect use?

A. Microsoft Power Platform admin center
B. Success by Design
C. The Business Value Toolkit
D. Microsoft Cloud Adoption Framework for Azure

Correct Answer: C

Why: The Business Value Toolkit is purpose-built to translate solution metadata and telemetry into a quantified ROI/ROAI figure. The admin center provides operational monitoring, Success by Design is a delivery methodology, and CAF for Azure is an infrastructure adoption framework — none directly compute business value from telemetry.

Q2. SINGLE-SELECT MULTIPLE CHOICE

An agent ingests customer feedback from a cloud database and produces monthly sentiment reports that will directly influence marketing and product decisions. Before ingestion, what is the MOST important data-preparation step?

A. Sort the records alphabetically by customer name
B. Identify and remediate biased or unrepresentative data
C. Build a Power Automate flow to move the data
D. Convert the database to a flat CSV file

Correct Answer: B

Why: Because the output will directly shape business decisions, unaddressed bias in the source data would propagate into skewed sentiment conclusions. Sorting or file-format changes do not affect data quality or fairness; a Power Automate flow only moves data, it doesn't clean it.

Q3. SINGLE-SELECT MULTIPLE CHOICE

A team wants generative AI to answer natural-language questions using a governed subset of corporate data that already exists as a Power BI semantic model connected to Dataverse via DirectQuery. What should the architect do to make this data usable as a trustworthy grounding source?

A. Export the semantic model to a spreadsheet
B. Share the Power BI workspace with all users
C. Endorse the semantic model
D. Rebuild the model as a canvas app

Correct Answer: C

Why: Endorsing (certifying/promoting) the semantic model signals to AI grounding services and end users that the model is trusted and governed, making it a discoverable, high-confidence grounding source. Sharing the workspace controls access but not trust signaling; exporting or rebuilding as an app breaks the live DirectQuery connection.

Q4. SINGLE-SELECT MULTIPLE CHOICE

A manufacturer is evaluating an agent that automates supplier invoice processing and is specifically worried about budget overruns. The design must weigh total cost of ownership, expected automation savings, and whether to extend existing AI capability. What should the architecture include?

A. A break-even analysis only

B. Immediate adoption of a prebuilt agent regardless of fit
C. Training a new custom model from scratch by default
D. A return on AI investment (ROAI) analysis

Correct Answer: D

Why: ROAI analysis is the structured method that combines TCO, projected savings, and build-vs-extend trade-offs into a single investment decision — directly addressing the budget-overrun concern. A break-even view alone ignores ongoing TCO, and jumping straight to 'prebuilt' or 'train custom' skips the financial evaluation step entirely.

Q5. SINGLE-SELECT MULTIPLE CHOICE

A Copilot Studio agent must answer customer-support questions using data spread across a ticketing system, two ERP applications, and spreadsheets stored in a document library. What should the architect recommend to give the agent accurate, timely access to all four sources?

A. Hard-code static prompts containing sample answers
B. Implement Power Platform connectors to each source
C. Rebuild all four systems inside one Dataverse environment
D. Ask users to manually paste data into the chat

Correct Answer: B

Why: Connectors let the agent retrieve live, accurate data from each heterogeneous system without duplicating or migrating it. Static prompts go stale immediately; consolidating four production systems into Dataverse is disproportionate and disruptive; manual paste does not scale.

Q6. SINGLE-SELECT MULTIPLE CHOICE

A Copilot Studio agent grounded in Dataverse returns irrelevant or inaccurate answers. Investigation shows the underlying table has duplicate, missing, and inconsistently formatted values. What is the FIRST corrective action?

A. Retrain the agent's conversational topics
B. Add an adaptive card to the response
C. Verify and cleanse the Dataverse data
D. Add more example utterances to the agent

Correct Answer: C

Why: Grounding quality is capped by source data quality — cleansing duplicates, filling gaps, and standardizing formats addresses the root cause. Retraining topics, adaptive cards, and extra utterances only change the conversation layer and won't fix answers built on dirty data.

Q7. SINGLE-SELECT MULTIPLE CHOICE

A financial-services support team currently manually reviews transaction histories to flag potential fraud before escalating cases to human analysts. The architect must automate the screening step while guaranteeing a human still makes the final call. What should be recommended?

A. An autonomous agent that auto-closes non-fraud cases without review
B. Microsoft 365 Copilot in Word to finalize fraud policy documents
C. A task agent that generates fraud-risk scores for a human analyst to review
D. Exporting all data to a data lake for offline Power BI analysis

Correct Answer: C

Why: A task/insight agent that scores risk but stops short of closing cases preserves human-in-the-loop decision-making, matching the requirement. Auto-closing cases removes the human review the business explicitly requires; Word Copilot and a data lake export don't address the real-time screening need.

Q8. SINGLE-SELECT MULTIPLE CHOICE

A retailer wants a single, consistent product catalog available to both a Dynamics 365 Commerce Copilot agent and a separate Power Apps inventory app. What should the architect recommend?

A. Let each agent independently scrape product pages from SharePoint
B. Duplicate the catalog into a separate custom table per agent
C. Configure prompts that read directly from vendor PDF files
D. Centralize the product catalog in Dataverse and expose it to both agents

Correct Answer: D

Why: A single Dataverse-hosted catalog eliminates duplication and version drift, giving both consuming agents one governed source of truth. Scraping, duplicating per agent, and reading from unstructured vendor PDFs all introduce inconsistency and staleness.

Q9. SINGLE-SELECT MULTIPLE CHOICE

While reviewing a customer-support agent, an architect finds a topic instructed only to 'help the customer as best you can,' with no grounding source, alongside a second topic that correctly retrieves facts from a curated knowledge base. What should the architect do to the first topic to ensure consistent, accurate behavior?

A. Leave it unchanged since flexible wording improves conversational tone
B. Replace the vague instruction with specific goals, constraints, and a defined grounding/knowledge source
C. Delete the topic entirely and rely on the second topic only
D. Increase the model's creativity setting to compensate

Correct Answer: B

Why: Vague, ungrounded instructions produce inconsistent and potentially fabricated answers. The fix is the same discipline used for the second topic: specific goals/constraints and an authoritative grounding source — not simply raising creativity, which increases variability, or deleting functionality the business still needs.

Q10. SINGLE-SELECT MULTIPLE CHOICE

An architect is configuring a Copilot Studio in-app help agent for Dynamics 365 finance and operations apps. To minimize inaccurate ('hallucinated') answers while still allowing the agent to answer product questions, what should the architect do regarding knowledge sources?

A. Enable general AI knowledge and disable all custom knowledge sources
B. Enable both general AI knowledge and custom knowledge sources with equal priority
C. Upload curated custom knowledge sources and disable general AI (public) knowledge
D. Disable both general AI knowledge and custom knowledge sources

Correct Answer: C

Why: Restricting the agent to curated, uploaded custom knowledge sources and turning off open-ended general AI knowledge minimizes the risk of the model inventing plausible-sounding but incorrect product guidance — a standard mitigation for grounded enterprise help agents.

Part 2 — Multi-Select Multiple Choice (Q11–Q20)

Q11. MULTI-SELECT MULTIPLE CHOICE

Which factors should be included when calculating return on AI investment (ROAI) for a proposed agent deployment? (Select all that apply.)

A. Total cost of ownership (build, run, and maintain costs)
B. Expected efficiency savings compared to the current manual process
C. Whether to extend an existing AI capability versus building new
D. The office furniture budget for the project team
E. Adoption and usage telemetry once the agent is live

Correct Answer: A, B, C, E

Why: ROAI is built from TCO, projected savings versus the current process, a build-vs-extend decision, and post-launch adoption telemetry to validate the estimate. Office furniture cost is unrelated to AI investment value.

Q12. MULTI-SELECT MULTIPLE CHOICE

Which requirements point to a CUSTOM Copilot Studio agent rather than a prebuilt Dynamics 365 AI agent template? (Select all that apply.)

A. The agent must be accessible from within Microsoft Teams as a low-code solution
B. The agent must connect to other agents that are selected dynamically based on their description
C. The agent must use business logic stored outside the core application
D. The agent must exactly replicate an out-of-the-box Dynamics 365 template with no changes
E. The agent must be embedded and accessible directly inside Dynamics 365 Supply Chain Management

Correct Answer: A, B, C, E

Why: Teams-embedded low-code delivery, dynamic multi-agent orchestration by description, use of external business logic, and embedded accessibility inside a D365 app all require the flexibility of a custom Copilot Studio agent. Replicating a template unchanged is, by definition, the prebuilt-agent path.

Q13. MULTI-SELECT MULTIPLE CHOICE

A CISO wants to monitor an AI agent for security and compliance risk. Which signals are directly relevant to that objective? (Select all that apply.)

A. How much sensitive data was accessed during a given agent run
B. Which user accessed the sensitive data
C. The visual theme color used in the agent's chat window
D. Unusually high volumes of sensitive-data access by a single user
E. The number of PowerPoint decks created that quarter

Correct Answer: A, B, D

Why: Sensitive-data volume per run, the identity of the accessing user, and anomalous access spikes by a single user are the core signals for a security/compliance audit. UI color and unrelated document counts carry no security-risk signal.

Q14. MULTI-SELECT MULTIPLE CHOICE

Which metrics should a CEO track to assess whether AI agents adhere to Responsible AI practices on a quarterly basis? (Select all that apply.)

A. Reliability

B. Interpretability
C. Fairness
D. Compliance
E. Total number of Teams channels created

Correct Answer: A, B, C, D

Why: A Responsible AI quarterly review, using Microsoft-recommended tooling, should verify reliability, interpretability, fairness, and compliance. Teams channel counts are an adoption/collaboration metric, not a Responsible AI metric.

Q15. MULTI-SELECT MULTIPLE CHOICE

Which practices help ensure Dataverse-grounded agent responses remain accurate and relevant over time? (Select all that apply.)

A. Periodically verifying and cleansing the underlying Dataverse data
B. Adding representative example user inputs to agent training
C. Randomly deleting older records to keep the table small
D. Reviewing and refining topic/grounding configuration based on user feedback
E. Disabling all logging so users feel less monitored

Correct Answer: A, B, D

Why: Ongoing data cleansing, representative training examples, and feedback-driven refinement sustain grounding quality. Randomly deleting records risks losing needed data, and disabling logging removes the visibility required to detect and fix quality issues.

Q16. MULTI-SELECT MULTIPLE CHOICE

A company centralizes its product catalog in Dataverse so multiple agents can use it as a shared grounding source. Which benefits does this design deliver? (Select all that apply.)

A. A single source of truth reduces version drift between agents
B. Governance and security controls can be applied once, centrally
C. Each agent automatically gets a different, independently maintained copy
D. Update effort is reduced because catalog changes are made in one place
E. It eliminates the need for any data quality validation

Correct Answer: A, B, D

Why: Centralization delivers a single source of truth, centralized governance, and single-point maintenance. It does not create independent per-agent copies (that's the opposite, duplicated-table pattern) and does not eliminate the ongoing need for data quality checks.

Q17. MULTI-SELECT MULTIPLE CHOICE

Which executive responsibilities are typically part of an enterprise AI-adoption governance model? (Select all that apply.)

A. CTO selecting and prioritizing which agents to deploy first
B. CIO ensuring appropriate sensitivity/security labels are applied to agent data
C. CFO analyzing ROI of deployed agents
D. CISO discovering and inventorying AI resources for audit
E. Any employee may independently approve new AI agents for production without review

Correct Answer: A, B, C, D

Why: A defined AI governance model assigns agent prioritization to the CTO, data labeling assurance to the CIO, ROI analysis to the CFO, and resource discovery/audit to the CISO. Unrestricted, unreviewed production approval by any employee undermines the entire governance model.

Q18. MULTI-SELECT MULTIPLE CHOICE

Which conditions indicate that an organization should extend an EXISTING AI capability rather than build a new custom model? (Select all that apply.)

A. An existing model already performs the required task with acceptable accuracy
B. TCO analysis shows extension is materially cheaper than a full custom build
C. The business requirement is entirely novel with no comparable existing capability
D. Time-to-value is a critical constraint and a prebuilt/extensible option meets requirements
E. Governance mandates every solution must be custom-built regardless of fit

Correct Answer: A, B, D

Why: Extending makes sense when an existing capability already fits, is cheaper per TCO analysis, and speeds time-to-value. A genuinely novel requirement may justify custom build, and a blanket 'always custom' governance rule ignores cost/time trade-offs entirely.

Q19. MULTI-SELECT MULTIPLE CHOICE

Which items should be documented for every proposed AI agent before deployment, per a mature AI-adoption governance process? (Select all that apply.)

A. A detailed business use case
B. A defined target audience/persona
C. An estimated ROI comparing current process cost to new solution cost
D. The personal performance review of the project sponsor
E. Data sensitivity classification for the sources the agent will use

Correct Answer: A, B, C, E

Why: Mature AI governance requires a documented use case, audience profile, ROI comparison, and data sensitivity classification before deployment. A sponsor's personal performance review is unrelated to solution governance.

Q20. MULTI-SELECT MULTIPLE CHOICE

Which factors should drive the choice between a topic triggered by an exact trigger phrase and a topic selected by description-based matching in Copilot Studio? (Select all that apply.)

A. Whether interactions need to feel more natural/conversational (favors description-based matching)
B. Whether the business requires deterministic, exact-phrase activation for compliance-critical flows (favors trigger phrases)
C. The number of connectors configured in the environment
D. Whether topics must be selected from a broader library of dynamically composed agents (favors description-based matching)
E. The default currency setting of the environment

Correct Answer: A, B, D

Why: Description-based topic selection improves conversational flexibility and supports dynamic multi-agent composition, while exact trigger phrases suit deterministic, compliance-sensitive flows. Connector counts and currency settings do not influence this design decision.

Part 3 — Drag and Drop (Q21–Q40)

Q21. DRAG AND DROP

Match each Copilot Studio analytics metric to the executive role it primarily helps. (Match technology/data to scenario.)

Draggable items: CTO | CIO | CFO | CISO

Target	Correct item placed
User feedback ratings on response quality, to trigger reengineering discussions	CTO
Sensitivity-label coverage across data sources used by agents	CIO
Agent usage estimator output compared to actual measured efficiency (ROI)	CFO
Volume of sensitive data accessed per run and by which user	CISO

Why: The CTO owns response-quality feedback that triggers reengineering; the CIO owns sensitivity labeling; the CFO compares estimated vs. actual ROI using usage data; the CISO monitors sensitive-data access volume and identity for audit and risk detection.

Q22. DRAG AND DROP

Order the recommended steps for planning an AI-powered business solution from first to last.

Draggable items: Pilot the solution with a limited audience | Define the business use case and target audience | Calculate estimated ROI/ROAI | Scale to full production with governance controls | Select a candidate business process

Target	Correct item placed
Step 1	Select a candidate business process
Step 2	Define the business use case and target audience
Step 3	Calculate estimated ROI/ROAI
Step 4	Pilot the solution with a limited audience
Step 5	Scale to full production with governance controls

Why: Sound planning starts by identifying a candidate process, defining the use case/audience, quantifying expected ROI, validating with a limited pilot, and only then scaling into governed production — skipping ahead risks unvalidated, ungoverned deployments.

Q23. DRAG AND DROP

Match each grounding data preparation technology/action to the correct scenario.

Draggable items: Endorse the semantic model | Cleanse and de-duplicate the Dataverse table | Apply a connector | Centralize into a shared Dataverse table

Target	Correct item placed
A Power BI DirectQuery model must be signaled as trusted for AI grounding	Endorse the semantic model
An agent returns irrelevant answers traced to dirty source data	Cleanse and de-duplicate the Dataverse table
An agent needs live access to a ticketing system it doesn't own	Apply a connector
Two agents need one consistent product catalog	Centralize into a shared Dataverse table

Why: Each grounding problem maps to a distinct fix: certification/endorsement for trust signaling, data cleansing for quality issues, connectors for live external access, and centralization for multi-agent consistency.

Q24. DRAG AND DROP

Order the steps for validating a Copilot Studio prompt/topic before production release.

Draggable items: Define specific goals, constraints, and expected tone | Attach a curated grounding/knowledge source | Test with representative and edge-case user inputs | Identify vague or ungrounded instructions | Approve for production release

Target	Correct item placed
Step 1	Identify vague or ungrounded instructions
Step 2	Define specific goals, constraints, and expected tone
Step 3	Attach a curated grounding/knowledge source
Step 4	Test with representative and edge-case user inputs
Step 5	Approve for production release

Why: Prompt validation starts by finding vague instructions, then tightening goals/constraints, attaching authoritative grounding, testing broadly, and only then approving — approving before testing or grounding risks inconsistent production behavior.

Q25. DRAG AND DROP

Match each agent type to the business requirement it best satisfies.

Draggable items: Autonomous agent | Task/insight agent | Prebuilt Dynamics 365 agent | Custom Copilot Studio agent

Target	Correct item placed
Runs event-triggered workflows in the background without a user prompt	Autonomous agent
Generates a risk score for a human analyst to review before action	Task/insight agent
Ships out-of-the-box for a supported Dynamics 365 scenario	Prebuilt Dynamics 365 agent
Must be embedded in Teams and use business logic stored outside the core app	Custom Copilot Studio agent

Why: Autonomous agents run without a live user prompt; task/insight agents surface analysis for human decision-making; prebuilt agents address common D365 scenarios out of the box; custom Copilot Studio agents provide the flexibility for Teams embedding and external business logic.

Q26. DRAG AND DROP

Order the ROAI (return on AI investment) analysis workflow.

Draggable items: Measure actual adoption/efficiency post-launch | Estimate current-process cost baseline | Compare actual results to the original estimate | Estimate total cost of ownership of the AI solution | Calculate estimated ROI before deployment

Target	Correct item placed
Step 1	Estimate current-process cost baseline
Step 2	Estimate total cost of ownership of the AI solution
Step 3	Calculate estimated ROI before deployment

Step 4	Measure actual adoption/efficiency post-launch
Step 5	Compare actual results to the original estimate

Why: ROAI starts by baselining current-process cost, then estimating TCO of the AI solution, producing a pre-deployment ROI estimate, and after go-live measuring actual adoption/efficiency to compare against the original estimate — closing the value-realization loop.

Q27. DRAG AND DROP

Match each data source integration technology to the correct scenario for a multi-system customer-support agent.

Draggable items: Power Platform connector | Azure AI Search incremental indexing | Model router | Adaptive card

Target	Correct item placed
Keep a large external knowledge index fresh without full reprocessing	Azure AI Search incremental indexing
Give the agent live read access to ServiceNow, D365 F&SCM, and SharePoint Excel files	Power Platform connector
Direct different query types to the most appropriate underlying model	Model router
Render a structured, interactive response inside the chat window	Adaptive card

Why: Incremental indexing keeps a large index current efficiently; connectors give live multi-system access; a model router directs queries to the best-fit model; adaptive cards render structured interactive UI in a chat response.

Q28. DRAG AND DROP

Order the process for centralizing a shared grounding data source used by multiple agents.

Draggable items: Expose the table to the consuming agents | Identify duplicate/legacy copies across agents | Migrate authoritative data into one governed Dataverse table | Apply sensitivity labels and access controls | Decommission the duplicate copies

Target	Correct item placed
Step 1	Identify duplicate/legacy copies across agents
Step 2	Migrate authoritative data into one governed Dataverse table
Step 3	Apply sensitivity labels and access controls
Step 4	Expose the table to the consuming agents
Step 5	Decommission the duplicate copies

Why: Centralization requires first finding existing duplicates, consolidating into one governed table, securing it with labels/access controls, exposing it to consumers, and only then retiring the old duplicate copies — retiring too early breaks dependent agents.

Q29. DRAG AND DROP

Match each knowledge-source configuration choice to the outcome it is designed to minimize or maximize for a generative help agent.

Draggable items: Enable only curated custom knowledge sources | Enable general AI (public) knowledge | Enable both with equal weighting | Disable all knowledge sources

Target	Correct item placed
Minimizes hallucination risk for enterprise product guidance	Enable only curated custom knowledge sources

Maximizes breadth of general topic coverage but raises inaccuracy risk	Enable general AI (public) knowledge
Creates ambiguity about which source the agent should trust first	Enable both with equal weighting
Produces no useful answers at all	Disable all knowledge sources

Why: Restricting to curated sources minimizes hallucination; enabling public knowledge broadens coverage but raises inaccuracy risk; enabling both without priority creates trust ambiguity; disabling everything removes the agent's ability to answer at all.

Q30. DRAG AND DROP

Order the discovery process the CTO should follow before selecting a prebuilt Dynamics 365 AI agent.

Draggable items: Preview agent capabilities with business stakeholders | Review available prebuilt agent templates | Select the highest-value agent for the initial phase | Identify which business operations would benefit most | Confirm the agent aligns with the documented use case and ROI estimate

Target	Correct item placed
Step 1	Review available prebuilt agent templates
Step 2	Identify which business operations would benefit most
Step 3	Preview agent capabilities with business stakeholders
Step 4	Confirm the agent aligns with the documented use case and ROI esti
Step 5	Select the highest-value agent for the initial phase

Why: Discovery starts with reviewing available templates, mapping them to business operations that would benefit, previewing with stakeholders, validating alignment with the documented use case/ROI, and finally selecting the highest-value candidate for phase one.

Q31. DRAG AND DROP

Match each Responsible AI evaluation dimension to its correct definition.

Draggable items: Reliability | Interpretability | Fairness | Compliance

Target	Correct item placed
The system performs consistently and as intended across expected conditions	Reliability
Stakeholders can understand why the model produced a given output	Interpretability
The system treats different user/demographic groups equitably	Fairness
The system adheres to applicable regulatory and organizational policy requirements	Compliance

Why: These four dimensions are distinct: reliability is about consistent behavior, interpretability is about explainability, fairness is about equitable treatment, and compliance is about adherence to regulation/policy — all four are reviewed together in a quarterly RAI assessment.

Q32. DRAG AND DROP

Order the steps to remediate a Dataverse-grounded agent that is returning irrelevant answers due to poor data quality.

Draggable items: Re-test agent responses against known-good queries | Diagnose whether the root cause is data or configuration | Cleanse, de-duplicate, and standardize the Dataverse table | Document the fix and monitor ongoing quality | Review recent user feedback and failure examples

Target	Correct item placed

Step 1	Review recent user feedback and failure examples
Step 2	Diagnose whether the root cause is data or configuration
Step 3	Cleanse, de-duplicate, and standardize the Dataverse table
Step 4	Re-test agent responses against known-good queries
Step 5	Document the fix and monitor ongoing quality

Why: Effective remediation starts with reviewing real failures, diagnosing root cause, applying the data-quality fix, re-testing against known-good queries, and documenting/monitoring — jumping straight to a fix without diagnosis risks treating the wrong problem.

Q33. DRAG AND DROP

Match each stakeholder concern to the correct planning artifact required before an AI agent is approved for deployment.

Draggable items: CFO's ROI concern | CISO's audit concern | CEO's Responsible AI concern | CIO's data protection concern

Target	Correct item placed
Documented ROI comparison of current vs. proposed process cost	Documented ROI comparison of current vs. proposed process cost
AI resource inventory entry with owner and risk classification	AI resource inventory entry with owner and risk classification
Quarterly reliability, interpretability, fairness, and compliance assessment	Quarterly reliability, interpretability, fairness, and compliance assessment
Sensitivity label assignment on all data sources the agent uses	Sensitivity label assignment on all data sources the agent uses

Why: Each executive's governance concern is satisfied by a specific artifact: ROI documentation for the CFO, an audited inventory entry for the CISO, a quarterly RAI assessment for the CEO, and sensitivity labeling for the CIO.

Q34. DRAG AND DROP

Order the recommended sequence for extending an existing AI capability instead of building a new custom model.

Draggable items: Confirm the existing capability meets accuracy requirements | Assess TCO of extension vs. full custom build | Pilot the extended capability on the new use case | Identify a candidate existing AI capability | Approve extension based on ROAI comparison

Target	Correct item placed
Step 1	Identify a candidate existing AI capability
Step 2	Confirm the existing capability meets accuracy requirements
Step 3	Assess TCO of extension vs. full custom build
Step 4	Pilot the extended capability on the new use case
Step 5	Approve extension based on ROAI comparison

Why: The extend-vs-build decision should start with identifying a candidate capability, confirming it meets accuracy needs, comparing TCO against a custom build, piloting on the real use case, and finally approving based on the ROAI comparison.

Q35. DRAG AND DROP

Match each Copilot Studio topic-selection mechanism to the scenario it best supports.

Draggable items: Exact trigger phrase | Description-based (intent) matching | Autonomous event trigger | Manual escalation to a human

Target	Correct item placed
A compliance-critical workflow requires deterministic, unambiguous activation	Exact trigger phrase
The interaction should feel natural and conversational without memorized phrases	Description-based (intent) matching
The agent must act on a system event without a user typing anything	Autonomous event trigger
A fraud risk score exceeds the confidence threshold for automated action	Manual escalation to a human

Why: Exact trigger phrases suit deterministic compliance flows; description-based matching suits natural conversation; autonomous event triggers suit background automation; and manual escalation preserves human judgment when confidence is insufficient.

Q36. DRAG AND DROP

Order the steps for designing a task agent that screens transactions for fraud while keeping a human as the final decision-maker.

Draggable items: Route flagged cases to a human analyst for final decision | Ingest and score transaction data for fraud risk | Define the escalation threshold and required human sign-off | Log every automated decision for audit | Define the data sources and sensitivity classification

Target	Correct item placed
Step 1	Define the data sources and sensitivity classification
Step 2	Ingest and score transaction data for fraud risk
Step 3	Define the escalation threshold and required human sign-off
Step 4	Route flagged cases to a human analyst for final decision
Step 5	Log every automated decision for audit

Why: Design should start by classifying the data used, then scoring risk, defining the escalation/sign-off threshold, routing flagged items to a human, and logging every decision for audit — preserving human-in-the-loop control throughout.

Q37. DRAG AND DROP

Match each business-value evidence type to the executive that primarily relies on it.

Draggable items: CFO | CTO | CISO | CEO

Target	Correct item placed
Usage estimator output compared against measured efficiency gains	CFO
User satisfaction feedback trends that may trigger reengineering	CTO
Monthly compliance and security audit findings	CISO
Quarterly Responsible AI scorecard across all deployed agents	CEO

Why: The CFO reviews the usage estimator vs. actual efficiency; the CTO tracks satisfaction trends to decide on reengineering; the CISO conducts monthly audits; the CEO reviews the quarterly Responsible AI scorecard across the portfolio.

Q38. DRAG AND DROP

Order the steps for reducing abandoned agent conversations flagged by a CFO as an inefficient use of Copilot Studio credits.

Draggable items: Compare abandonment rate to resolved-conversation rate over time | Identify the topics/steps where users most often abandon | Redesign the affected topics for clarity and shorter paths | Re-measure abandonment after the redesign | Establish a baseline abandoned-vs-resolved metric

Target	Correct item placed
Step 1	Establish a baseline abandoned-vs-resolved metric
Step 2	Compare abandonment rate to resolved-conversation rate over time
Step 3	Identify the topics/steps where users most often abandon
Step 4	Redesign the affected topics for clarity and shorter paths
Step 5	Re-measure abandonment after the redesign

Why: Effective remediation starts with a baseline, trends the ratio over time, pinpoints the specific abandonment points, redesigns those topics, and re-measures to confirm improvement — skipping the baseline makes it impossible to prove the redesign helped.

Q39. DRAG AND DROP

Match each MCP (Model Context Protocol) usage scenario to the correct architectural role.

Draggable items: Existing R&D; MCP server | Custom Copilot Studio agent | Compliance data endpoint | Product specification query

Target	Correct item placed
Provides structured tool/data access that the agent can call for grounding answers	Existing R&D MCP server
The consumer that invokes the MCP server to answer manufacturing questions	Custom Copilot Studio agent
A sensitive data category requiring strict access boundaries within the MCP server	Compliance data endpoint
The specific business question the agent must answer using MCP-exposed data	Product specification query

Why: The MCP server exposes structured tool/data access; the custom agent is the consumer calling it; compliance data within the server needs tightly scoped access; and the product specification query is the business intent the whole chain exists to satisfy.

Q40. DRAG AND DROP

Order the lifecycle for a prebuilt Dynamics 365 Supply Chain Management AI agent from discovery through initial adoption.

Draggable items: Deploy to a limited pilot audience | Discover available agent templates | Preview capabilities with operations stakeholders | Prioritize the highest-value template | Measure early adoption and feedback

Target	Correct item placed
Step 1	Discover available agent templates
Step 2	Preview capabilities with operations stakeholders
Step 3	Prioritize the highest-value template
Step 4	Deploy to a limited pilot audience
Step 5	Measure early adoption and feedback

Why: The adoption lifecycle proceeds from discovering templates, previewing them with the stakeholders who will use them, prioritizing the highest-value option, piloting with a limited audience, and measuring adoption/feedback before wider rollout.

Part 4 — Drop-Down: Complete the Architecture (Q41–Q60)

Q41. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the grounding architecture for a generative help-and-guidance agent in Dynamics 365 finance and operations apps.

Blank	Drop-down options	Correct selection
Custom knowledge sources	Must be uploaded to the agent / Must never be configured / Are optional and ignored by the model	Must be uploaded to the agent
AI general (public) knowledge	Must be enabled for broader coverage / Must be disabled for the agent / Has no effect on accuracy	Must be disabled for the agent

Why: To minimize inaccurate responses in an enterprise help agent, curated custom knowledge sources must be uploaded and trusted, while open-ended general AI knowledge should be disabled so the model cannot fabricate answers from ungoverned public information.

Q42. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the executive responsibility matrix for AI adoption governance.

Blank	Drop-down options	Correct selection
CTO	Select and prioritize agents to deploy / Assign sensitivity labels / Approve marketing budget	Select and prioritize agents to deploy
CIO	Analyze ROI / Assign appropriate sensitivity/security labels to agent data / Discover and inventory AI resources	Assign appropriate sensitivity/security labels to agent data
CISO	Discover and inventory AI resources for audit / Select agents to deploy / Finalize the product roadmap	Discover and inventory AI resources for audit

Why: Each executive owns a distinct governance lane: the CTO prioritizes deployment, the CIO ensures correct data labeling, and the CISO maintains a discoverable, auditable inventory of AI resources.

Q43. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the architecture for making a Power BI semantic model usable as an AI grounding source.

Blank	Drop-down options	Correct selection
Trust signal step	Endorse the semantic model / Rename the workspace / Delete the DirectQuery connection	Endorse the semantic model
Connection type retained	DirectQuery to Dataverse / A one-time CSV export / A screenshot of the report	DirectQuery to Dataverse

Why: Endorsing the model signals trust to AI grounding tooling, while keeping the live DirectQuery connection to Dataverse ensures the grounding source stays current rather than becoming a stale, one-time export.

Q44. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the remediation architecture for an agent that returns inaccurate answers due to poor Dataverse data quality.

Blank	Drop-down options	Correct selection
Root-cause action	Verify and cleanse the Dataverse data / Increase the model's temperature / Add more decorative UI elements	Verify and cleanse the Dataverse data
Secondary reinforcement	Add representative example user inputs to training / Delete the agent / Disable grounding entirely	Add representative example user inputs to training

Why: The primary fix targets the data quality issue at its source; adding representative training examples further reinforces correct behavior once the underlying data is trustworthy.

Q45. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the multi-source customer-support agent architecture.

Blank	Drop-down options	Correct selection
Live access to ERP + ticketing + SharePoint files	Power Platform connectors / Static hard-coded prompts / Manual data re-entry	Power Platform connectors
Keeping a large external search index current	Incremental indexing in Azure AI Search / Full manual reindex weekly by an admin / Disable the search index	Incremental indexing in Azure AI Search

Why: Connectors give live, governed access to disparate systems, while incremental indexing efficiently keeps a large external search index fresh without a costly full reprocessing cycle.

Q46. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the financial evaluation architecture for a supplier-invoice automation agent focused on avoiding budget overruns.

Blank	Drop-down options	Correct selection
Primary financial analysis	Return on AI investment (ROAI) analysis / Break-even analysis only / No financial analysis needed	Return on AI investment (ROAI) analysis
Build-vs-extend decision input	TCO comparison against extending existing AI capability / Marketing team's brand guidelines / UI color palette	TCO comparison against extending existing AI capabil

Why: A full ROAI analysis, informed by a TCO-based build-vs-extend comparison, directly targets the budget-overrun concern by quantifying both cost and expected savings before committing.

Q47. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the fraud-screening automation architecture that must keep a human as final decision-maker.

Blank	Drop-down options	Correct selection
Automated component	Task agent generating fraud-risk scores / Autonomous agent auto-closing all cases / Word Copilot finalizing poli	Task agent generating fraud-risk scores
Human-in-the-loop step	Escalation to a human analyst for final decision / No review required / Automatic email to the customer only	Escalation to a human analyst for final decision

Why: A task agent surfaces risk scores without taking final action, and every flagged case must still be escalated to a human analyst — preserving the required human decision point.

Q48. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the product-catalog architecture shared across a Commerce Copilot agent and a Power Apps inventory app.

Blank	Drop-down options	Correct selection
Storage location	Centralized Dataverse table / Per-agent duplicate custom tables / Vendor PDF files	Centralized Dataverse table
Exposure method	Expose the shared table to both agents / Recreate the catalog manually in each app / Email the catalog to each	Expose the shared table to both agents

Why: A single, centrally governed Dataverse table exposed to both consumers avoids duplication, version drift, and manual synchronization overhead.

Q49. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the topic-design architecture for consistent, accurate customer-support responses.

Blank	Drop-down options	Correct selection
Vague topic fix	Add specific goals, constraints, and a grounding source / Increase response length only / Rename the topic	Add specific goals, constraints, and a grounding sourc
Knowledge-retrieval topic	Keep the existing curated knowledge base grounding / Remove all grounding for consistency / Switch to genera	Keep the existing curated knowledge base grounding

Why: An ungrounded, vague topic needs explicit goals/constraints and a real grounding source, while a topic that already retrieves from a curated knowledge base should keep that grounding rather than be weakened toward open-ended public knowledge.

Q50. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the Responsible AI quarterly assessment architecture required by the CEO.

Blank	Drop-down options	Correct selection
Tooling requirement	Microsoft-recommended RAI assessment tooling / Any thi[illegible]	Microsoft-recommended RAI assessment tooling
Required dimensions	Reliability, interpretability, fairness, compliance / Speed, c[illegible]	Reliability, interpretability, fairness, compliance

Why: The CEO's quarterly review specifically requires Microsoft-recommended tooling verifying reliability, interpretability, fairness, and compliance — general performance metrics like latency or cost alone do not satisfy a Responsible AI review.

Q51. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the conversation-efficiency architecture the CFO uses to detect wasted Copilot Studio credits.

Blank	Drop-down options	Correct selection
Key comparison metric	Abandoned vs. resolved conversations per day / Total nu[illegible]	Abandoned vs. resolved conversations per day
Follow-up action if aban[illegible]	Redesign the affected topics for clarity / Immediately shut[illegible]	Redesign the affected topics for clarity

Why: Comparing abandoned to resolved conversations reveals inefficient credit usage; a high abandonment rate should trigger topic redesign rather than an abrupt shutdown or being ignored.

Q52. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the sensitive-data audit architecture required by the CISO.

Blank	Drop-down options	Correct selection
Volume metric	Amount of sensitive data accessed per agent run / Numb[illegible]	Amount of sensitive data accessed per agent run
Identity metric	Which user accessed the sensitive data / The agent's dis[illegible]	Which user accessed the sensitive data

Why: A security audit needs both how much sensitive data was touched per run and who accessed it — an unusually high volume tied to a single user is the specific risk signal the CISO is watching for.

Q53. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the custom-agent architecture required to answer questions using R&D; product specifications maintained outside the core application.

Blank	Drop-down options	Correct selection
Integration mechanism	Existing Model Context Protocol (MCP) server / Manual c[illegible]	Existing Model Context Protocol (MCP) server
Access approach	Reuse the existing MCP server via structured tool access[illegible]	Reuse the existing MCP server via structured tool access

Why: Since R&D; already exposes product specifications and compliance data through an existing MCP server, the agent should reuse that structured tool/data access rather than duplicating sensitive proprietary data elsewhere.

Q54. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the topic-selection architecture for a more conversational custom agent.

Blank	Drop-down options	Correct selection
Selection mechanism	Description-based (intent) matching / Exact trigger phrase[illegible]	Description-based (intent) matching
Benefit delivered	More natural, flexible interactions / Strict deterministic co[illegible]	More natural, flexible interactions

Why: Selecting topics by description/intent rather than a fixed trigger phrase produces more natural conversational behavior, at the cost of the strict determinism that exact-phrase triggers provide for compliance-critical flows.

Q55. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the multi-agent orchestration architecture where the custom agent must connect to other agents chosen dynamically.

Blank	Drop-down options	Correct selection
Selection basis for other agents	Each agent's description / A fixed hard-coded list only / Alphabetical order of agent names	Each agent's description
Resulting orchestration style	Dynamic, description-driven orchestration / Static, single-agent-only design / Manual human routing for every re	Dynamic, description-driven orchestration

Why: Selecting connected agents based on their description (rather than a static list) enables dynamic orchestration that can scale as new agents are added, without requiring the base agent to be reconfigured each time.

Q56. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the ROI validation architecture the CFO uses each quarter.

Blank	Drop-down options	Correct selection
Estimation tool used pre-deployment	Copilot Studio agent usage estimator / Power BI Desktop only / Azure Cost Management only	Copilot Studio agent usage estimator
Post-deployment comparison	Estimated ROI vs. actual measured efficiencies and adoption / Estimated ROI vs. competitor pricing / Estimated	Estimated ROI vs. actual measured efficiencies and adoption

Why: The CFO uses the usage estimator for the pre-deployment projection, then compares that estimate against actual measured efficiency and adoption after go-live to validate whether the investment delivered as expected.

Q57. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the data-preparation architecture for an agent that will drive marketing and product decisions from customer sentiment.

Blank	Drop-down options	Correct selection
Primary risk to mitigate	Biased or unrepresentative source data / Slow network latency / Incorrect font choice	Biased or unrepresentative source data
Mitigation step	Identify and address the bias before ingestion / Sort records alphabetically / Convert file format only	Identify and address the bias before ingestion

Why: Because the output directly informs business decisions, the key data-preparation risk is bias in the source data, which must be identified and remediated before ingestion — cosmetic changes like sorting or reformatting do not address this risk.

Q58. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the discovery architecture the CTO follows for prebuilt Dynamics 365 Supply Chain Management agent selection.

Blank	Drop-down options	Correct selection
Discovery input	Available prebuilt agent templates / Vendor sales brochures only / Competitor product reviews	Available prebuilt agent templates
Validation step before selection	Preview agent capabilities with relevant stakeholders / Skip preview and deploy directly to production / Ask only	Preview agent capabilities with relevant stakeholders

Why: The CTO reviews available templates and then previews the highest-priority candidate's capabilities with the relevant business operation stakeholders before committing to deployment — skipping stakeholder validation risks selecting a poor-fit agent.

Q59. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the escalation-quality architecture the CTO uses to decide when to trigger a reengineering discussion.

Blank	Drop-down options	Correct selection
Signal monitored	User feedback on the quality of agent responses / Number of licenses purchased / Server CPU utilization only	User feedback on the quality of agent responses
Trigger condition	Consistently poor feedback over time / A single one-off negative comment / Any feedback at all, positive or nega	Consistently poor feedback over time

Why: An escalated reengineering discussion should be triggered by a consistent pattern of poor quality feedback, not by an isolated negative comment or by feedback volume alone regardless of sentiment.

Q60. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the extend-vs-build decision architecture for a new AI capability.

Blank	Drop-down options	Correct selection
First evaluation step	Assess whether an existing AI capability already fits the need / Immediately begin training a new model / Purcha	Assess whether an existing AI capability already fits th
Deciding financial input	TCO/ROAI comparison of extend vs. build / Marketing team preference / Office relocation budget	TCO/ROAI comparison of extend vs. build

Why: The decision should begin by checking fit against existing capability, then be resolved using a TCO/ROAI-based financial comparison between extending that capability and building new — not by defaulting to a custom build or unrelated budget considerations.

Part 5 — Case Studies (Q61–Q70)

Case Study 1 — Fabrikam Precision Manufacturing

Scenario: Fabrikam Precision Manufacturing designs and produces patented industrial components using Dynamics 365 Finance, Dynamics 365 Supply Chain Management, and Dynamics 365 Commerce across several legal entities and Azure subscriptions. Executive AI-adoption responsibilities are: CTO — select one prebuilt Dynamics 365 Supply Chain Management agent and one custom Copilot Studio agent for the initial rollout; CIO — ensure correct sensitivity labels on all data used by agents; CFO — analyze ROI of deployed agents using the Copilot Studio agent usage estimator; CISO — discover and inventory all AI resources monthly for audit; CEO — ensure quarterly Responsible AI review of reliability, interpretability, fairness, and compliance. Engineering already operates a Model Context Protocol (MCP) server exposing product specifications and compliance data. The planned custom agent must: run as a low-code solution inside Microsoft Teams, be embedded directly in Dynamics 365 Supply Chain Management, answer manufacturing questions using existing product-specification data, select connected agents dynamically based on description (not a fixed trigger phrase), and reuse business logic stored outside the core application.

Q61. SINGLE-SELECT MULTIPLE CHOICE

Which action should the CTO take FIRST before selecting the prebuilt Dynamics 365 Supply Chain Management agent?

A. Deploy the agent directly to all production users
B. Review available agent templates and preview capabilities with relevant business operation stakeholders
C. Ask the CFO to select the agent instead
D. Skip the prebuilt agent and build everything custom

Correct Answer: B

Why: Per the scenario, the CTO must first discover available templates and preview the highest-priority option with stakeholders before committing — direct production deployment or bypassing the prebuilt option entirely skips the required discovery step.

Q62. MULTI-SELECT MULTIPLE CHOICE

Which of the custom agent's stated requirements specifically justify building it in Copilot Studio rather than using only a prebuilt template? (Select all that apply.)

A. Must run as a low-code Teams solution
B. Must dynamically select connected agents by description
C. Must reuse business logic stored outside the core application
D. Must exactly copy an existing prebuilt template with zero modification
E. Must be embedded directly inside Dynamics 365 Supply Chain Management

Correct Answer: A, B, C, E

Why: Teams embedding, dynamic description-based agent selection, external business-logic reuse, and direct D365 SCM embedding all require the flexibility of a custom Copilot Studio build. Exact replication of a prebuilt template with no change is, by definition, the prebuilt path — not a justification for custom development.

Q63. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the knowledge-source configuration for the custom agent answering product-specification questions.

Blank	Drop-down options	Correct selection
Primary data access mechanism	Reuse the existing engineering MCP server / Manually copy [illegible] product specifica	Reuse the existing engineering MCP server

Compliance data handling	Apply strict, scoped access controls within the MCP integration	Apply strict, scoped access controls within the MCP in / Expose all compliance data openly to any user

Why: The agent should reuse the existing MCP server rather than duplicating proprietary specification data, while compliance data — being especially sensitive given Fabrikam's patents — requires scoped access controls consistent with least privilege.

Q64. DRAG AND DROP

Match each Fabrikam executive to the AI governance artifact they are accountable for.

Draggable items: CFO | CISO | CEO | CIO

Target	Correct item placed
ROI analysis using the Copilot Studio agent usage estimator	CFO
Monthly AI resource inventory for audit	CISO
Quarterly Responsible AI review (reliability, interpretability, fairness, compliance)	CEO
Sensitivity labels on data used by agents	CIO

Why: Each artifact maps directly to the executive named in the scenario: the CFO owns ROI/usage-estimator analysis, the CISO owns the monthly inventory/audit, the CEO owns the quarterly Responsible AI review, and the CIO owns sensitivity labeling.

Q65. SINGLE-SELECT MULTIPLE CHOICE

The custom agent must connect to other agents that are selected based on their description rather than a fixed trigger phrase. What is the primary architectural benefit of this design choice for Fabrikam?

A. It guarantees the lowest possible Azure compute cost
B. It allows the orchestration to scale as new agents are added, without reconfiguring trigger phrases each time
C. It removes the need for any sensitivity labeling
D. It eliminates the need for a CISO audit

Correct Answer: B

Why: Description-based agent selection lets the multi-agent orchestration grow dynamically as Fabrikam adds new agents over time, rather than requiring every new agent to be manually wired to a fixed trigger phrase — it has no bearing on cost, labeling, or audit requirements.

Case Study 2 — Northwind Health Group

Scenario: Northwind Health Group operates regional clinics and a claims-processing division. Claims staff currently spend significant manual effort reviewing claim histories for potential billing anomalies before escalating suspicious claims to a senior reviewer. Leadership wants to introduce AI to reduce manual effort while keeping a human as the final decision-maker on any denied or flagged claim. Patient and claims data is highly regulated and must remain within an approved geographic region. Northwind already has a certified, endorsed Power BI semantic model (DirectQuery to Dataverse) summarizing claims trends, and a separate, ungoverned spreadsheet-based claims extract used informally by one regional team. Leadership also wants a monthly comparison between the number of AI-assisted conversations that are abandoned versus resolved, to judge whether licensing spend is being used efficiently.

Q66. SINGLE-SELECT MULTIPLE CHOICE

Which automation design correctly reduces manual claims-review effort while preserving the required human-in-the-loop control?

A. An autonomous agent that automatically denies claims flagged as anomalous

B. A task agent that generates an anomaly/risk score and routes flagged claims to a senior reviewer for final decision
C. A generative writing agent that drafts claim denial letters for immediate mailing
D. Removing the review step entirely to save time

Correct Answer: B

Why: A task agent that scores risk and routes flagged items to a human reviewer directly satisfies Northwind's requirement to reduce manual effort while keeping a human as the final decision-maker; auto-denial, auto-mailing letters, or removing review entirely all violate that requirement.

Q67. MULTI-SELECT MULTIPLE CHOICE

Which data-governance actions should be taken regarding Northwind's two claims data sources before using them for AI grounding? (Select all that apply.)

A. Continue using the endorsed semantic model as a trusted grounding source
B. Migrate the ungoverned spreadsheet extract into a governed, sensitivity-labeled Dataverse table before use
C. Use the ungoverned spreadsheet as-is because it is faster to access
D. Confirm claims data processing stays within the approved geographic region
E. Disable all logging on claims-related agent interactions

Correct Answer: A, B, D

Why: The already-endorsed semantic model is a trusted source and can continue to be used; the ungoverned spreadsheet must be brought under governance (labeled, centralized) before use; and regional data-residency must be explicitly confirmed given the regulated data. Using an ungoverned source as-is or disabling logging both violate governance and auditability requirements.

Q68. DROP-DOWN (ARCHITECTURE COMPLETION)

Complete the data residency and grounding architecture for Northwind's claims AI solution.

Blank	Drop-down options	Correct selection
Regulated claims data processing region	Must stay within the approved geographic region / May be processed in any region for convenience / Region do	Must stay within the approved geographic region
Preferred grounding source	The endorsed semantic model via DirectQuery / The informal spreadsheet extract / An unrelated public dataset	The endorsed semantic model via DirectQuery

Why: Regulated claims data must remain within the approved region, and the already-endorsed, live-connected semantic model is the correct grounding source over the ungoverned spreadsheet or an unrelated public dataset.

Q69. SINGLE-SELECT MULTIPLE CHOICE

Leadership wants a monthly comparison of abandoned versus resolved AI-assisted conversations to evaluate licensing efficiency. What should be established FIRST to make this comparison meaningful over time?

A. A baseline abandoned-vs-resolved ratio measured before any changes are made
B. An immediate topic redesign with no measurement
C. A decision to cancel the AI licenses pre-emptively
D. A survey of unrelated employee satisfaction

Correct Answer: A

Why: Without an initial baseline, leadership has no way to judge whether the ratio is improving, worsening, or stable month over month — redesigning topics or cancelling licenses without first establishing a baseline would be premature and unmeasurable.

Q70. SINGLE-SELECT MULTIPLE CHOICE

Which financial evaluation approach should Northwind use to justify the claims-review automation investment to its board?

A. A qualitative opinion from the claims team only

B. A return on AI investment (ROAI) analysis comparing current manual-review cost to the automated solution's TCO and projected savings

C. A comparison of competitor marketing materials

D. No financial justification is needed for internal tools

Correct Answer: B

Why: A structured ROAI analysis — comparing the current manual process cost against the automated solution's total cost of ownership and projected savings — gives the board a defensible, quantified basis for the investment decision, unlike an informal opinion or no analysis at all.

Part 6 — Hotspot / Yes-No Case Scenarios (Q71–Q80)

Q71-73. HOTSPOT / YES-NO (CASE SCENARIO)

Scenario: A CTO reviews a newly proposed custom Copilot Studio agent for a manufacturing client. The agent will be embedded in Microsoft Teams, will select which of several connected agents to invoke based on the description of each agent's purpose, and will pull grounding data from an existing MCP server that has not yet had sensitivity labels applied to its compliance dataset.

Statement	Yes / No
Selecting connected agents by description rather than a fixed trigger phrase is an appropriate design choice for a conversational, low-cod	Yes
Because the MCP server already exists and is trusted internally, sensitivity labeling of the compliance dataset can be safely skipped befo	No
The agent's Teams embedding and low-code delivery model, on their own, are sufficient justification to skip a documented business use	No

Why: Description-based agent selection fits the conversational, dynamically-composed design goal. However, an existing internal MCP server being 'trusted' does not remove the requirement to apply sensitivity labels to compliance data before production use — governance requirements apply regardless of the data source's origin. Similarly, a low-code or Teams-embedded delivery model does not exempt the agent from the standard governance requirement for a documented use case and ROI estimate.

Q74-76. HOTSPOT / YES-NO (CASE SCENARIO)

Scenario: A CFO is reviewing quarterly analytics for three deployed Copilot Studio agents. Agent A shows a rising trend of abandoned conversations relative to resolved ones. Agent B shows actual measured efficiency closely matching the original usage-estimator projection. Agent C has had no ROI review since its initial deployment eighteen months ago.

Statement	Yes / No
Agent A's rising abandonment trend, on its own, is sufficient evidence to conclude the agent should be immediately decommissioned with	No
Agent B's alignment between actual efficiency and the original usage-estimator projection indicates the initial ROAI estimate was well-cal	Yes
Agent C going eighteen months without an ROI review is consistent with the recommended practice of periodic, ongoing value-realization	No

Why: A rising abandonment trend should trigger investigation and possible topic redesign, not immediate decommissioning without root-cause analysis. Close alignment between actual and estimated efficiency is exactly the signal that confirms a well-calibrated ROAI estimate. Going 18 months without any ROI review contradicts the expectation of periodic quarterly (or otherwise ongoing) value-realization comparisons.

Q77-80. HOTSPOT / YES-NO (CASE SCENARIO)

Scenario: A CISO's monthly audit of a claims-processing agent shows: (1) one user account accessed an unusually large volume of sensitive claims data in a single day, (2) all agent interactions are logged centrally with user identity captured, (3) the agent's grounding data includes an ungoverned spreadsheet extract that has never been reviewed for data quality or sensitivity, and (4) the agent's autonomous workflow can approve low-value claims without any human review step.

Statement	Yes / No
The unusually large single-user access volume is a signal that should be investigated as a potential security risk, not dismissed as routin	Yes
Centralized logging with captured user identity supports the CISO's ability to audit sensitive-data access as required.	Yes
Using the ungoverned, unreviewed spreadsheet extract as a grounding source is an acceptable practice as long as the agent's answers s	No

An autonomous workflow that approves low-value claims without any human review is consistent with best practice for claims-processing	No

Why: A single-user spike in sensitive-data access is precisely the anomaly pattern the CISO's monitoring is designed to catch and must be investigated. Centralized logging with identity capture is the correct control supporting auditability. An ungoverned, unreviewed data source should never be used for grounding regardless of how plausible its outputs appear — data quality and sensitivity review must happen first. Finally, even low-value claim approvals should retain a human-in-the-loop or clearly bounded, audited automation — fully autonomous approval without any review contradicts the human-in-the-loop requirement established for claims automation.