

Quick gut-check on your example question

The Copilot Studio knowledge-source question (Azure AI Search / Data Lake / Cosmos DB / Translator) is a **red flag pattern**, not a good template. The AB-100 source material never actually discusses which Azure data service backs a Copilot Studio knowledge source — that's Power Platform/Copilot Studio product trivia, not architecture-design content taught in these units. If you use this style verbatim, you're testing outside the syllabus, and a sharp candidate will correctly argue there's no textual basis for "the" answer.

3 red flags for this exam bank

1. **Out-of-syllabus product trivia.** The docs teach *process and governance frameworks* (ALM gates, RACI, RAI principles), not specific Azure SKU trivia. Any question requiring "which specific Azure service" knowledge not stated in the text is ungrounded and indefensible in a dispute.
2. **Redundant near-duplicate units.** "Design security agents for Microsoft clouds," "Design governance models for AI agents," and "Design model security for responsible AI" restate the same RBAC/DLP/observability/threat-protection content three times with slightly different framing. This makes single-best-answer questions risky — two units may each imply a "correct" but differently-worded answer.
3. **Source document errors baked into visuals.** The Copilot "test maturity ladder" diagram literally has a duplicated "Level 4" entry (two different Level 4s, no Level 5-in-order integrity), and the closing loop text has a stray trailing asterisk ("Test environment.*"). If a drag-and-drop ordering question is built directly off that diagram, the answer key itself is broken.

Strong vs. weak sections

Strong (good for tough, defensible questions):

- ALM promotion gates (Foundry, custom models, D365 F&SC) — clean, enumerated, unambiguous gate criteria. Great for single-select and dropdown.
- Responsible AI six principles — clean taxonomy, perfect for multi-select and matching.
- Audit trail requirements (model vs. data vs. Foundry control-plane) — clearly separated, testable categories.
- Data ALM phases A–G — sequential, well-labeled, ideal for drag-and-drop ordering.

Weak (avoid building hard questions here without fixing source content first):

- Copilot test case blueprint table — the PDF extraction is garbled/truncated (columns cut off mid-sentence), so field-matching questions risk citing broken source text.
- Test maturity ladder — has the duplicate Level 4 defect noted above.
- The three security/governance units — too much overlap to safely test "which unit says X" without ambiguity.

Now building the full bank as a file (60+ questions across all formats) so it's easy to review and reuse. Two things worth flagging before you ship this to candidates:

1. **MS-06 and MS-08** each have 4 correct answers, not 3 — fix the option count or relabel them "select all that apply."
2. **Don't build ordering/drag-drop items off the Copilot Test Maturity Ladder graphic** — the source PDF has a duplicated "Level 4" entry, so the sequence itself is broken in the original material.

Everything else in the bank is directly traceable to specific tables, gate criteria, or named lists in your two source docs, so you've got a defensible answer key if a candidate disputes a question.

AB-100 Exam Question Bank

Architect AI Solutions for Business Productivity — Solution Architect Level

Prepared as an internal exam-setter deliverable for jitendraco.

Grounded in the two source modules: *Manage testing AI-powered business solutions* and *Design ALM process for AI-powered business solutions* + *Design responsible AI security, governance, risk management, and compliance*.

0. Exam-Setter Analysis (Summary)

Top 5 Strongest Questions in This Bank (with reasoning)

#	Question (short form)	Why it's strong
1	SS-03 (below): Which environment gate criterion applies at Gate D→E in the AI data ALM process	Directly traceable to a specific named phase table in the source; only one option matches the literal gate text — no ambiguity.

2	MS-02: Select all controls that belong to "defense in depth" for agent security	Source explicitly lists identity, data governance, observability, and threat protection as the layered stack — clean multi-select with a clear boundary between correct/incorrect options.
3	DD-04: Order the Data ALM phases A→G	Phases are explicitly labeled A through G in sequence in the source — unambiguous correct order, tests real process comprehension not memorization of trivia.
4	CS-A Q3 (case study): Given a canary rollout requirement, which ALM gate must the team satisfy before Prod	Forces the candidate to apply Gate D→E logic to a scenario rather than recall a list — genuine architect-level reasoning.
5	YN-01 Q2: "A model can be promoted to Prod without a documented rollback plan."	Tests a hard governance rule (Gate 2 requires rollback plan) with a definitive Yes/No answer traceable to source text — no interpretation needed.

3 Red Flags (recap, see above for full explanation)

1. Out-of-syllabus product/SKU trivia not stated in source text.
2. Heavy redundancy across the three security/governance units risks "which unit is this from" ambiguity — avoid unit-attribution questions.
3. Source diagram defects (duplicate "Level 4" in test maturity ladder) — do not build ordering questions directly from that diagram without correcting it first.

Strong Sections

- ALM promotion gates (Foundry / custom models / D365 F&SC / Copilot Studio agents) — precise, enumerable criteria.
- Responsible AI's six principles — clean, closed taxonomy.
- Audit trail categories (model changes vs. data changes vs. Foundry control-plane) — well partitioned.
- Data ALM Phase A–G lifecycle — sequential and complete.
- RACI table for AI data — good source for matching/drag-drop.

Weak Sections

- Copilot test case blueprint table (garbled PDF extraction, columns cut off) — do not test exact field wording.

- Test maturity ladder graphic — contains a duplicate "Level 4" defect; unusable for ordering questions as-is.
 - The three overlapping security/governance units — good for concept questions, unsafe for "which module teaches this" questions.
-

1. Single-Select Multiple Choice (10) — "Choose the best architectural decision"

SS-01 A solution architect is designing an environment strategy for Copilot Studio agents, connectors, and actions. The customer wants to guarantee no untested change reaches end users directly. Which environment design principle should the architect enforce?

- A. Allow direct editing in Production for urgent fixes
- B. Use managed solutions only in Test and Prod, with no direct editing in Production
- C. Use a single shared environment for Dev, Test, and Prod to reduce cost
- D. Allow makers to publish directly from Dev to Prod

Answer: B — The source explicitly lists "No direct editing in production" and "Managed solutions only in Test and Prod" as core environment design principles for Copilot Studio ALM.

SS-02 An architect must decide how a Copilot-driven test case generation strategy is initiated. What should be defined *before* Copilot is used to generate test cases?

- A. The number of Copilot licenses required
- B. The testing objectives (test categories, business rules, risks, compliance requirements)
- C. The UI theme for the test dashboard
- D. The Azure region hosting the test environment

Answer: B — Section 2.1 "Define the testing objectives" is explicitly the required first step before prompting Copilot.

SS-03 In the AI data ALM process, which of the following is required as evidence at **Gate D → E (Stage & Approve → Deploy & Serve)**?

- A. A signed reproducibility log only

- B. CAB approval, deployment runbooks, and rollback plans
- C. A profiling report and lineage IDs
- D. A data contract and catalog policy proof

Answer: B — Per the Gate table, Gate D→E requires "CAB approval; deployment runbooks and rollback plans ready."

SS-04 A team is designing a promotion path for a Microsoft Foundry agent. Which activity belongs specifically to **Promotion Gate 2 (Test → Prod)**, not Gate 1 (Dev → Test)?

- A. Initial safety and guardrail checks
- B. Verification of data source mappings
- C. Human validation of agent reasoning
- D. Code or prompt quality review

Answer: C — "Human validation of agent reasoning" is listed only under Gate 2 (Test → Prod); the other three are Gate 1 criteria.

SS-05 An architect is designing identity for AI agents operating across Microsoft clouds. What is the recommended authentication approach for agent-to-Azure calls?

- A. Static API keys stored in the agent configuration
- B. Shared service principal credentials across all agents
- C. Managed identities to remove secrets and simplify rotation
- D. User-delegated OAuth tokens cached indefinitely

Answer: C — The source states "Prefer managed identities for agent-to-Azure authentication to remove secrets and simplify rotation."

SS-06 Which control is described as enforcing that agents only use data "in the right places, for the right duration"?

- A. Threat protection
- B. Data governance and protection (DLP, sensitivity labels, residency, retention)
- C. Observability and cost governance
- D. Development and interoperability standards

Answer: B — This is the stated goal of the "Data governance and protection" section.

SS-07 A financial services customer requires that AI-driven journal postings and purchase order approvals cannot occur without safeguards. Which control category from the D365 Finance and Supply Chain ALM unit directly addresses this?

- A. Version control for models
- B. Guardrails for business-critical processes (e.g., posting journals, approving purchase orders)
- C. Synthetic dataset generation
- D. Environment variables for endpoints

Answer: B — Explicitly named under "Risk & Compliance Controls" in the D365 F&SC ALM unit.

SS-08 Which of the following best distinguishes the "red" dataset from the "gold" dataset in the environment strategy for AI data ALM?

- A. Red is used only in Production; gold is used only in Dev
- B. Red is mutable and experimental; gold is frozen and promoted
- C. Red contains PII; gold never contains PII
- D. Red and gold refer to model versions, not datasets

Answer: B — Directly stated: "Red gold datasets pattern: red (mutable, experimental) vs. gold (frozen, promoted)."

SS-09 An architect reviewing an AI solution for Responsible AI adherence needs to evaluate whether the system "empowers people of all abilities and backgrounds." Which RAI principle does this map to?

- A. Fairness
- B. Transparency
- C. Inclusiveness
- D. Accountability

Answer: C — Inclusiveness is defined exactly this way in the RAI principles list.

SS-10 When validating data residency for a Copilot Studio agent, what should the architect confirm regarding unpublished agents and preview features?

- A. They always share identical residency rules with published agents
- B. They may follow different residency rules than published agents and must be validated separately
- C. They are exempt from all residency requirements
- D. They automatically inherit tenant-wide DLP policy with no further review needed

Answer: B — The source explicitly instructs architects to validate "Whether unpublished agents and preview features follow different residency rules."

2. Multi-Select Multiple Choice (10) — "Select two or three correct answers"

MS-01 (Select 3) Which of the following are explicitly listed as AI assets that must be governed under ALM for Copilot Studio?

- A. Agents (conversational or autonomous)
- B. Custom connectors
- C. Actions, skills, and prompt assets
- D. The end user's personal calendar
- E. Employee performance reviews

Answer: A, B, C

MS-02 (Select 3) Per the "defense in depth" approach to AI agent security, which controls form part of the layered posture?

- A. Identity and access boundaries
- B. Secure data governance
- C. Disabling all logging to reduce cost
- D. Monitoring and anomaly detection
- E. Granting broad access "just in case" a future task needs it

Answer: A, B, D

MS-03 (Select 2) Which two elements are required as evidence for **Gate 2 (Test → Prod)** in the custom AI model ALM lifecycle?

- A. Security and compliance approvals granted
- B. Rollback plan confirmed
- C. Training data meets quality standards
- D. Documentation created

Answer: A, B (C and D belong to Gate 1: Dev → Test)

MS-04 (Select 3) Which of the following are named Responsible AI principles in the source material?

- A. Fairness
- B. Reliability and Safety
- C. Scalability
- D. Privacy and Security
- E. Portability

Answer: A, B, D

MS-05 (Select 2) Which two practices are recommended to reduce AI model drift and data poisoning risk?

- A. Track data lineage for all training sources
- B. Disable retraining permanently once a model reaches production
- C. Validate newly introduced data before use in retraining pipelines
- D. Allow unrestricted public internet ingestion for freshness

Answer: A, C

MS-06 (Select 3) Which categories of AI vulnerabilities are explicitly described in the "Analyze AI vulnerabilities" unit?

- A. Prompt manipulation risks

- B. Model behavior vulnerabilities
- C. Data exposure vulnerabilities
- D. Network cable degradation
- E. Identity, access, and RBAC gaps

Answer: A, B, C, E — *(Note: this is actually 4 correct; use as a "select all that apply" variant, or trim option E to force exactly 3 — flagged for calibration.)*

MS-07 (Select 2) Which two controls are recommended specifically for securing grounding data retrieval workflows (e.g., RAG)?

- A. Connector-level authorization restricting which data types the model can query
- B. Allowing the model unrestricted field access to improve recall
- C. Structured query filtering preventing access to disallowed fields
- D. Removing all logging to reduce retrieval latency

Answer: A, C

MS-08 (Select 3) Which of the following are listed as recommended environments in the ALM strategy for AI in Dynamics 365 Customer Service / Customer Engagement?

- A. Development (DEV)
- B. Test/Validation (TEST)
- C. Sandbox-of-Sandboxes (nested sandbox)
- D. Pre-Production (UAT/PREPROD)
- E. Production (PROD)

Answer: A, B, D, E — *(4 correct — same calibration note as MS-06; consider trimming distractors to force exactly 3.)*

MS-09 (Select 2) According to the audit trail unit, which two attributes should recommended audit logs include?

- A. Immutable, timestamped change records
- B. Full raw content of every prompt and dataset for convenience

- C. Role-based attribution linked to identity provider
- D. Automatic deletion after 24 hours regardless of workload risk

Answer: A, C — (B is explicitly discouraged: "logs capture metadata, not content, to avoid unnecessary exposure.")

MS-10 (Select 3) Which of these are named as AI assets to govern in the D365 Finance and Supply Chain ALM unit?

- A. AI models (prediction, scoring, classification, anomaly detection)
- B. Prompts and instructions for Copilot assistance
- C. Knowledge sources (documents, task guides, structured data entities)
- D. Employee salary bands
- E. Marketing campaign creative assets

Answer: A, B, C

3. Drag-and-Drop (10)

DD-01 — Match the ALM environment to its purpose (D365 F&SC) Match: DEV / TEST / PROD Targets:

1. Execute approved AI capabilities in live financial and supply chain workloads
2. Build and iterate AI models, prompts, orchestration logic, and integrations
3. Validate with safe, anonymized production-like data; perform regression checks

Answer: DEV→2, TEST→3, PROD→1

DD-02 — Order the Action lifecycle (Copilot Studio ALM) Items: Build, Design, Monitor, Promote, Validate

Answer (correct order): Design → Build → Validate → Promote → Monitor

DD-03 — Match the RAI principle to its definition Match: Fairness / Transparency / Accountability / Privacy and Security Targets:

1. Solutions should be understandable, with clear disclosures on how AI is used

2. AI systems should treat all groups equitably
3. Organizations retain responsibility for decisions made by AI
4. Protect personal and organizational data through strong controls

Answer: Fairness→2, Transparency→1, Accountability→3, Privacy and Security→4

DD-04 — Order the AI Data ALM Phases Items: Deploy & Serve, Plan & Catalog, Develop & Evaluate, Ingest & Prepare, Stage & Approve, Operate & Monitor, Evolve & Retire

Answer (correct order = A→G): Plan & Catalog → Ingest & Prepare → Develop & Evaluate → Stage & Approve → Deploy & Serve → Operate & Monitor → Evolve & Retire

DD-05 — Match the control area to its Gate in the AI Data ALM Gate table Match: A→B / B→C / C→D / D→E / E→F Targets:

1. Evaluation & safety
2. Compliance & residency
3. Catalog & ownership
4. Runtime readiness
5. Data quality & lineage

Answer: A→B→3, B→C→5, C→D→1, D→E→2, E→F→4

DD-06 — Match the connector release flow step to its position Items: Apply DLP, Approve Security, Author, Publish to Prod, Validate Auth

Answer (correct order): Author → Validate Auth → Apply DLP → Approve Security → Publish to Prod

DD-07 — Match the RACI role to its accountability for Residency Configuration (AI data RACI table) Match: Data Owner / AI Architect / Security-Compliance / Platform Admin Targets: Consulted / Accountable / Responsible / Responsible

Answer: Data Owner→Consulted (C), AI Architect→Accountable (A), Security/Compliance→Responsible (R), Platform Admin→Responsible (R)

DD-08 — Order the Responsible AI Lifecycle stages Items: Build, Deploy, Design, Monitor, Validate

Answer (correct order): Design → Build → Validate → Deploy → Monitor

DD-09 — Match the vulnerability category to its example Match: Prompt manipulation / Data exposure / Identity & RBAC gaps / Agent workflow vulnerability
Targets:

1. Excessive permissions allow the model to access data it doesn't need
2. "Ignore previous instructions..." style override attempts
3. Autonomous tool use without proper guardrails
4. Attacker interacts with the model using elevated privileges

Answer: Prompt manipulation→2, Data exposure→1, Identity & RBAC gaps→4, Agent workflow vulnerability→3

DD-10 — Order the Continuous Improvement loop (Copilot Studio ALM) Items: Analyze, Improve, Monitor, Release, Validate

Answer (correct order): Monitor → Analyze → Improve → Release → Validate → (loops back to Monitor)

4. Dropdown / Complete-the-Diagram (10)

DO-01 "In the Environment & Data Flow diagram, data moves from Dev (Red data) → Test (Repro runs, evaluation sets) → Pre-Prod (___) → Prod (Gold only)." Options: [Gold candidates | Raw ingestion | Draft prompts | Archived snapshots]

Answer: Gold candidates

DO-02 "In the Residency Decision Tree, if in-region capacity is unavailable and overflow is allowed for the workload tier, the architect should ____." Options: [Block the feature entirely | Enable cross-region processing under admin control | Delete the workload | Migrate the entire tenant]

Answer: Enable cross-region processing under admin control

DO-03 "In the RBAC design table, the Publisher role is scoped with ____ at the Azure Roles level." Options: [Least privilege | Full Owner access | No access | Global Admin]

Answer: Least privilege

DO-04 "In the Data Governance Layers diagram, the layer that sits between 'Agent Runtime' and 'User Output & Logging' is ____." Options: [DLP & Sensitivity Filters | Marketing Automation | Payment Processing | Identity Federation]

Answer: DLP & Sensitivity Filters

DO-05 "In the Retrieval Access Flow for grounding data, the correct order is: Prompt → ____ → Search Index → Sanitization Layer → Model Context Injection." Options: [Policy Check | Billing Check | UI Rendering | Backup Snapshot]

Answer: Policy Check

DO-06 "In the Model Hardening Blueprint, Step 2 ('Private Endpoints') focuses on ____." Options: [Using private network endpoints to limit exposure and control access | Detecting prompt injection | Validating training data lineage | Setting retention policy]

Answer: Using private network endpoints to limit exposure and control access

DO-07 "On the Copilot Studio environment diagram, DEV is (Unmanaged), TEST is (Managed), and PROD is (____)." Options: [Managed | Unmanaged | Hybrid | Sandboxed]

Answer: Managed

DO-08 "In the AI Data ALM Gate table, the evidence required at Gate B→C is ____." Options: [Profiling report; lineage IDs | CAB approval | Dashboards; alarms; budget guard | Residency mapping; approval memo]

Answer: Profiling report; lineage IDs

DO-09 "In the Copilot test case strategy map, after 'Copilot Test Case Drafts,' the next stage is ____." Options: [Architect Validation Pass | Final Test Case Library | Business Requirements | Architect Prompt Plan]

Answer: Architect Validation Pass

DO-10 "In the D365 F&SC environment table, the environment used to 'Validate with safe, anonymized production-like data' is ____." Options: [TEST | DEV | PROD | Sandbox-Preview]

Answer: TEST

5. Case Studies (2–3 scenarios, 8 questions total)

Case Study A — Contoso Retail Customer Service Copilot (3 questions)

Scenario: Contoso is deploying a Copilot Studio agent for Dynamics 365 Customer Service to summarize cases and suggest next-best actions. The agent will access customer interaction data, case history, and knowledge articles. Compliance requires that no direct edits happen in Production and that all prompts undergo regression testing before release. The security team also requires that sensitive customer data never appears in agent prompts unnecessarily.

A-Q1. Which environment should the team use to run regression tests for prompts, summarization consistency, and classification accuracy before Pre-Production?

- A. Dev
- B. Test/Validation
- C. Prod
- D. Pre-Prod only

Answer: B (Test/Validation) — regression testing for prompts and classification accuracy is explicitly assigned to the Test/Validation environment.

A-Q2. Which principle should guide how much customer data is exposed to the model in prompts?

- A. Maximize context by including full case history in every prompt
- B. Data minimization — use the minimum data required for the model's purpose
- C. Include raw PII to improve summarization accuracy
- D. Data does not need governance once inside a prompt

Answer: B — "Data minimization: Use the minimum data required for the model's purpose" is an explicit model security control.

A-Q3. Before deployment to Prod, which artifact confirms the release is safe to ship per the AI Release Readiness Checklist for D365 AI?

- A. A checklist confirming data quality, knowledge source alignment, prompt regression testing, safety/compliance verification, environment variables, deployment pipeline success, and monitoring dashboards readiness
- B. A single sign-off email from the product owner
- C. A verbal confirmation from the Dev team lead
- D. Nothing — Copilot Studio auto-validates releases

Answer: A — directly matches the "AI Release Readiness Checklist" items listed in the source.

Case Study B — Fabrikam Finance AI Forecasting Model (3 questions)

Scenario: Fabrikam is building a custom AI model to predict cash flow anomalies for Dynamics 365 Finance. The model must pass through Dev → Test → Prod, comply with data residency rules for financial data, and be monitored for prediction drift after go-live.

B-Q1. Which gate criteria must be satisfied before this model can move from Test to Prod?

- A. Only that training data meets quality standards
- B. Model evaluation thresholds met, no bias/unsafe output, performance verified under ERP workload constraints, security/compliance/residency reviews complete, and deployment package approved
- C. A single stakeholder sign-off with no formal evaluation
- D. Completion of Dev-stage prompt drafting only

Answer: B — this matches Gate 2 (TEST → PROD) criteria in the D365 F&SC ALM unit exactly.

B-Q2. After go-live, what should the team monitor to detect whether the model's predictions are degrading over time?

- A. UI click-through rate only
- B. Prediction drift (accuracy changes over time) and data drift (input patterns deviating from training data)
- C. Number of licenses consumed
- D. Screen resolution compatibility

Answer: B — explicitly named under "Operational Monitoring" in the D365 F&SC unit.

B-Q3. Which control specifically protects financial data residency requirements for this model?

- A. Sensitivity labelling enforcement and data residency requirements for financial data
- B. Reducing model parameter count
- C. Increasing token limits
- D. Disabling audit logging to reduce noise

Answer: A — listed directly under "Risk & Compliance Controls."

Case Study C — Northwind Traders Security Review (2 questions)

Scenario: Northwind's security team is reviewing an autonomous AI agent that reads support tickets and updates a knowledge base automatically. The review must assess identity design, data boundaries, and incident response readiness.

C-Q1. What is the recommended way to assign the agent's identity across Dev, Test, and Prod?

- A. Share a single identity across all environments for simplicity
- B. Assign a unique cloud identity per agent, per environment (prod, pre-prod, dev), with recorded ownership, version, and lifecycle metadata
- C. Use the security architect's personal credentials
- D. No identity is required for read-only agents

Answer: B — directly stated under "Agent identity."

C-Q2. As part of incident response readiness, what should be predefined for this agent?

- A. Steps to disable a compromised model endpoint, preserve logs, notify stakeholders, and recover safely
- B. A plan to permanently delete all logs immediately after any incident
- C. A policy that agents can never be disabled once published
- D. Nothing formal — ad hoc response is sufficient for read-only agents

Answer: A — matches "Incident response for AI models" guidance exactly.

6. Yes/No Hotspot Questions (cannot be revisited after submission)

Hotspot Scenario 1 — Copilot Studio Agent Governance

A company deploys a Copilot Studio agent that reads ticket data and updates knowledge articles automatically. Read each statement and select Yes or No.

#	Statement	Answer	Reason
1	The agent should be granted broad, standing access to all knowledge sources "in case it needs them later."	No	Violates least-privilege / data minimization principles stated throughout the security and governance units.
2	Publishing changes to Production requires an approval workflow, not direct editing.	Yes	"No direct editing in production" and "Require approvals for publishing to production" are explicit environment design principles.
3	Once deployed, the agent requires no further evaluation because it passed initial testing.	No	Source requires "Scheduled reviews for accuracy, data freshness, and risk reassessment" as ongoing lifecycle governance — agents are never "set and forget."

Hotspot Scenario 2 — Custom AI Model Promotion

An architect is preparing to promote a custom AI model for cash-flow anomaly detection from Test to Production.

#	Statement	Answer	Reason
1	A model can be promoted to Prod without a documented rollback plan if the evaluation metrics are strong.	No	Gate 2 (Test→Prod) explicitly requires "Rollback plan confirmed" as a mandatory criterion, independent of evaluation results.
2	Security and compliance approvals are required before promotion.	Yes	Explicitly listed under Gate 2 (Test → Prod) criteria.
3	Training on production knowledge sources is acceptable during the Develop & Evaluate phase to improve accuracy.	No	The source states "Train/iterate using Dev/Test data; never train on production knowledge."

Hotspot Scenario 3 — Data Residency and Compliance

A regulated customer wants to validate that a generative AI feature does not move data outside its approved region.

#	Statement	Answer	Reason
1	All Copilot Studio agents follow identical data residency rules regardless of publish status.	No	Source explicitly instructs validating "whether unpublished agents and preview features follow different residency rules."
2	DLP rules can be used to prevent sensitive data from being used in AI prompts or outputs.	Yes	Explicitly stated under Purview controls for Microsoft 365 Copilot.
3	In a regulated scenario, cross-region overflow processing should be enabled by default for performance.	No	Source states: "In regulated scenarios, set the default to in-region and require explicit approval to enable overflow processing."

Calibration Notes for the Exam Committee

- **MS-06 and MS-08** currently have 4 correct answers against a "select 3" instruction — either expand to "select all that apply" or trim one distractor before publishing.
- Avoid sourcing any drag-and-drop/ordering item from the **Copilot Test Maturity Ladder** graphic until the duplicate "Level 4" defect in the source PDF is corrected — currently unusable for a scored item.
- Recommend a technical review pass on the Copilot test case blueprint table (garbled column text in source) before building any exact-field-recall questions from it.

One thing to flag on this pass: **DD-01** (the Copilot Studio vs. Foundry agent deployment question) lands on the *same* answer for both agents, unlike your example where the two agents likely get differentiated methods. That's because the four requirements you gave rule out manual deployment equally for both — the source docs don't give a reason to split them. If you want a genuinely differentiated drag-and-drop like your example, add one requirement that's agent-specific (e.g., "the Foundry agent must deploy via infrastructure-as-code" vs. "the Copilot Studio agent must deploy via managed solution packages") — otherwise this item will read as a trick with only one real answer dressed up as two blanks.

AB-100 Exam Question Bank — v2 (Exam-Format Match)

Architect AI Solutions for Business Productivity

This version mirrors the official exam phrasing/format you supplied (scenario stem, NOTE lines, drag targets, case study structure, non-revisitable hotspot). Every answer is traceable to the two source modules — flagged where the source is thin.

1. Single-Select Multiple Choice (10)

"Choose the best architectural decision from four options"

SS-01 A solution architect is designing a Copilot Studio agent for a finance team. The agent must generate test cases automatically as business requirements change, and the test cases must follow a consistent, predictable structure every time. You need to recommend what the architect should provide to Copilot before requesting test cases. What should you include in the solution?

- A. A list of the finance team's Azure AD group memberships
- B. A test case blueprint template (Test ID, Purpose, Preconditions, Steps, Expected Results, Edge Cases, Dependencies)
- C. A Power BI dashboard of historical defect counts
- D. The Azure subscription ID for the Finance environment

Answer: B — "Providing this template in the prompt ensures Copilot outputs predictable, high-quality test cases."

SS-02 A company is promoting a Microsoft Foundry agent from Test to Production. The compliance team requires proof that a human has reviewed the agent's reasoning quality before go-live. What should you include as a Promotion Gate 2 requirement?

- A. Initial safety and guardrail checks
- B. Verification of data source mappings
- C. Human validation of agent reasoning
- D. Action and connector configuration review

Answer: C — Listed explicitly under Promotion Gate 2 (Test → Prod), not Gate 1.

SS-03 An architect is designing the environment strategy for a custom AI model used in fraud detection. The model must never be trained on live customer data. What should the architect enforce during the Model Development phase?

- A. Train exclusively using Dev/Test data; never train on production knowledge
- B. Use production data with encryption applied post-training
- C. Use a synthetic copy of production data refreshed nightly into Prod
- D. Allow the model to fine-tune itself continuously in Production

Answer: A — "Train/iterate using Dev/Test data; never train on production knowledge."

SS-04 A retail company wants centralized visibility into every prompt, tool call, and safety event generated by its AI agents, across all environments. What should you recommend?

- A. Per-agent local log files reviewed manually each week
- B. Centralized/unified logging that captures prompts, tool calls, errors, and safety events in a single workspace
- C. Disabling logging in Dev to reduce noise
- D. Storing logs only for agents flagged as high-risk

Answer: B — "Centralize telemetry for prompts, tool calls, errors, and safety events in a single workspace."

SS-05 A healthcare customer requires that AI-generated outputs never disclose data outside the tenant's approved geographic region, even during model inference or evaluation. What should the architect validate first?

- A. Where user prompts, context, and model inputs are processed
- B. The number of active Copilot licenses
- C. The color scheme of the agent's chat interface
- D. The marketing budget for the AI rollout

Answer: A — This is the first residency question the source instructs architects to answer.

SS-06 An architect must reduce the risk that an AI agent's tool-calling behavior triggers an unintended action against a production system. What should you recommend?

- A. Grant the agent's identity full Contributor rights across all environments
- B. Ensure the agent acts only within clearly defined tool capability boundaries, with auditing and rollback enabled
- C. Remove approval workflows to speed up automation
- D. Allow the agent to call any endpoint without an allow-list

Answer: B — Matches the agent/workflow-level vulnerability mitigation guidance directly.

SS-07 A company is evaluating whether a dataset is ready for promotion to Production in its AI ALM process. Which dataset should you select?

- A. The latest Dev extract with broad sample coverage
- B. A Test dataset that looks balanced but has partial lineage
- C. An immutable, versioned dataset with sensitivity labels and full lineage that passed evaluation gates
- D. Any dataset that delivers acceptable latency in a pilot

Answer: C — This is the correct answer from the source's own module-assessment question.

SS-08 An architect needs to reduce the blast radius if an AI agent is compromised by an adversarial prompt. What should you recommend as part of pre-production assurance?

- A. Skip red teaming to accelerate the release
- B. Run red team evaluations pre-production and after major changes; gate releases on findings closure
- C. Rely only on the model's built-in safety training
- D. Perform red teaming only once, at initial launch

Answer: B

SS-09 A company wants a single Model Card produced for every custom AI model prior to approval. What should the Model Card contain?

- A. Metrics, constraints, and recommended usage boundaries
- B. The developer's performance review
- C. The Azure invoice for the training run
- D. The end user's support ticket history

Answer: A — "Produce a Model Card containing metrics, constraints, and recommended usage boundaries."

SS-10 An architect is designing access control for a Retrieval-Augmented Generation (RAG) pipeline used by a customer support agent. What should you include to prevent the model from retrieving fields it is not authorized to access?

- A. Structured query filtering that prevents access to disallowed fields
- B. A larger context window
- C. A higher token limit per session
- D. Disabling the search index cache

Answer: A

2. Multi-Select Multiple Choice (10)

"Choose two or three correct answers." NOTE: Each correct selection is worth one point.

MS-01 Which two activities are required as evidence at **Gate D → E** in the AI data ALM process? Each correct answer presents part of the solution.

- A. CAB approval
- B. A signed reproducibility log
- C. Deployment runbooks and rollback plans
- D. A profiling report

Answer: A, C

MS-02 An architect is designing a defense-in-depth posture for enterprise AI agents. Which two controls belong in that layered stack? Each correct answer presents part of the solution.

- A. Identity and access boundaries

- B. Disabling logging to reduce operational cost
- C. Secure data governance
- D. Granting broad access so the agent can retrieve any data it may need later

Answer: A, C

MS-03 Which two criteria must be satisfied at **Gate 2 (Test → Prod)** for a custom AI model? Each correct answer presents part of the solution.

- A. Training data meets quality standards
- B. Security and compliance approvals granted
- C. Rollback plan confirmed
- D. Documentation created

Answer: B, C (*A and D belong to Gate 1: Dev → Test*)

MS-04 Which three are named Responsible AI principles in the source material? NOTE: Each correct selection is worth one point.

- A. Fairness
- B. Scalability
- C. Reliability and Safety
- D. Portability
- E. Privacy and Security

Answer: A, C, E

MS-05 Which two practices reduce the risk of model drift and data poisoning? Each correct answer presents part of the solution.

- A. Track data lineage for all training sources
- B. Permanently disable retraining after go-live
- C. Validate newly introduced data before use in retraining pipelines
- D. Allow unrestricted public internet ingestion for freshness

Answer: A, C

MS-06 Which two components should the security team evaluate when reviewing identity design for an autonomous agent? Each correct answer presents part of the solution.

- A. The agent's average response length
- B. Whether the agent uses a managed identity per environment
- C. The agent's marketing tagline
- D. Whether least-privilege role assignments are scoped narrowly

Answer: B, D

MS-07 Which three environments are recommended for the ALM strategy for AI in Dynamics 365 Customer Service? NOTE: Each correct selection is worth one point.

- A. Development (DEV)
- B. Sandbox-of-Sandboxes
- C. Test/Validation (TEST)
- D. Pre-Production (UAT/PREPROD)
- E. Personal Developer Environment

Answer: A, C, D

MS-08 Which two attributes should audit logs for AI model changes include? Each correct answer presents part of the solution.

- A. Full raw content of every prompt for convenience
- B. Immutable, timestamped change records
- C. Role-based attribution linked to identity provider
- D. Automatic deletion after 24 hours regardless of risk level

Answer: B, C

MS-09 Which two elements are part of the Retrieval Access Flow for grounding data? Each correct answer presents part of the solution.

- A. Policy Check
- B. Billing Reconciliation
- C. Sanitization Layer
- D. UI Theme Selection

Answer: A, C

MS-10 Which two are explicitly listed as AI assets to govern under ALM for AI in Dynamics 365 Finance and Supply Chain? Each correct answer presents part of the solution.

- A. AI models (prediction, scoring, classification, anomaly detection)
- B. Employee salary bands
- C. Prompts and instructions for Copilot assistance
- D. Marketing creative assets

Answer: A, C

3. Drag-and-Drop (10)

Matching technologies/methods to scenarios, or ordering implementation steps — mirrors the deployment-method example format.

DD-01 A company plans to implement an AI solution that will contain a Microsoft Copilot Studio agent and a Microsoft Foundry agent. The solution will be stored in a source code repository. You need to recommend a deployment method for each agent. The solution must meet the following requirements:

- A test environment must be used before a deployment to production.
- Production must be isolated from development and testing.
- The deployment must be repeatable and fully automated.
- The solution must **NOT** require manual intervention.

Which deployment method should you recommend for each agent? (Drag the appropriate method to each agent. Each method may be used once, more than once, or not at all.)

Methods: Manual solution export/import | Automated release pipeline (e.g., Azure DevOps/GitHub Actions) promoting managed solution packages through Dev → Test →

Prod | Ad hoc email of connection strings | One-time manual configuration in Production only

Targets: Copilot Studio agent | Microsoft Foundry agent

Answer:

- Copilot Studio agent → **Automated release pipeline promoting managed solution packages through Dev → Test → Prod**
 - Microsoft Foundry agent → **Automated release pipeline promoting managed solution packages through Dev → Test → Prod** *(Both agents require the same automated, repeatable, non-manual pattern; manual export/import fails the "no manual intervention" requirement for both.)*
-

DD-02 Order the Action lifecycle for Copilot Studio ALM. **Items:** Build, Design, Monitor, Promote, Validate

Answer: Design → Build → Validate → Promote → Monitor

DD-03 Match the connector release flow step to its correct position. **Items:** Apply DLP, Approve Security, Author, Publish to Prod, Validate Auth

Answer: Author → Validate Auth → Apply DLP → Approve Security → Publish to Prod

DD-04 A company needs to promote a dataset through its AI data ALM lifecycle. Match each dataset state to the correct environment. **Items:** Red (mutable, experimental) | Gold (frozen, promoted) **Targets:** Dev/Test | Pre-Prod/Prod

Answer: Red → Dev/Test; Gold → Pre-Prod/Prod

DD-05 Order the seven phases of the AI Data ALM process (A→G). **Items:** Deploy & Serve, Plan & Catalog, Develop & Evaluate, Ingest & Prepare, Stage & Approve, Operate & Monitor, Evolve & Retire

Answer: Plan & Catalog → Ingest & Prepare → Develop & Evaluate → Stage & Approve → Deploy & Serve → Operate & Monitor → Evolve & Retire

DD-06 A company is deploying AI features across Dynamics 365 Customer Service. Match each deployment target to its release mechanism. **Items:** PROD deployment |

PRE-PROD staging **Targets:** "Export managed solution packages, stage in PRE-PROD for final checks" | "Deploy to PROD using approved release pipelines (Azure DevOps or GitHub Actions)"

Answer: PRE-PROD staging → Export managed solution packages, stage for final checks; PROD deployment → Deploy using approved release pipelines

DD-07 Match the RAI Lifecycle stage to its correct order position. **Items:** Build, Deploy, Design, Monitor, Validate

Answer: Design → Build → Validate → Deploy → Monitor

DD-08 Match the RACI role to its accountability level for "Residency configuration" in the AI Data RACI table. **Items:** Data Owner, AI Architect, Security/Compliance, Platform Admin **Targets:** Consulted (C), Accountable (A), Responsible (R), Responsible (R)

Answer: Data Owner → C, AI Architect → A, Security/Compliance → R, Platform Admin → R

DD-09 A company must promote a Microsoft Foundry agent bundle. Match each governance activity to its correct promotion gate. **Items:** Verification of data source mappings | Regression test completion | Human validation of agent reasoning | Initial safety and guardrail checks **Targets:** Gate 1 (Dev→Test) | Gate 2 (Test→Prod)

Answer: Verification of data source mappings → Gate 1; Initial safety and guardrail checks → Gate 1; Regression test completion → Gate 2; Human validation of agent reasoning → Gate 2

DD-10 Order the Continuous Improvement loop used in Copilot Studio ALM. **Items:** Analyze, Improve, Monitor, Release, Validate

Answer: Monitor → Analyze → Improve → Release → Validate

4. Drop-Down / Diagram Completion (10)

Complete architectural/testing diagrams by selecting the correct component from a drop-down — mirrors the Dynamics 365 Sales/Finance testing example.

DO-01 You are designing end-to-end test scenarios for a business solution that uses Microsoft Dynamics 365 Sales and Dynamics 365 Finance. You need to ensure the solution meets the following test requirements:

- Properly exchanges data between the Dynamics 365 apps
- Aligns with defined user workflows and business processes

Select the correct testing type for each requirement from the drop-down list. Each testing type may be used once, more than once, or not at all.

Drop-down list: Unit testing | Integration/end-to-end testing | Regression testing | Scenario-based testing | Prompt evaluation

Requirement 1 — "Properly exchanges data between the Dynamics 365 apps": [drop-down] **Requirement 2 — "Aligns with defined user workflows and business processes":** [drop-down]

Answer:

- Requirement 1 → **Integration/end-to-end testing** — "end-to-end test scenarios ... validate cross-app data flow across multiple Dynamics 365 applications ... because AI output quality depends on consistent, trusted, and well-timed input data from across integrated systems."
 - Requirement 2 → **Scenario-based testing**
-

DO-02 Complete the Environment & Data Flow diagram: "Dev (Red data) → Test (Repro runs, evaluation sets) → Pre-Prod (___) → Prod (Gold only)"

Drop-down: Gold candidates | Raw ingestion | Draft prompts | Archived snapshots

Answer: Gold candidates

DO-03 Complete the Residency Decision Tree: "In-region capacity available? → No → Is overflow allowed for this workload tier? → Yes → ___"

Drop-down: Enable cross-region processing under admin control | Block the feature | Delete the workload | Migrate the tenant

Answer: Enable cross-region processing under admin control

DO-04 Complete the RBAC Design table: "Publisher role is scoped with ___ at the Azure Roles level."

Drop-down: Least privilege | Full Owner access | No access | Global Admin

Answer: Least privilege

DO-05 Complete the Data Governance Layers diagram: "Data Sources → Agent Runtime → ____ → User Output & Logging"

Drop-down: DLP & Sensitivity Filters | Marketing Automation | Payment Processing | Identity Federation

Answer: DLP & Sensitivity Filters

DO-06 Complete the Retrieval Access Flow: "Prompt → ____ → Search Index → Sanitization Layer → Model Context Injection"

Drop-down: Policy Check | Billing Check | UI Rendering | Backup Snapshot

Answer: Policy Check

DO-07 Complete the Model Hardening Blueprint: "Step 2, 'Private Endpoints,' focuses on ____."

Drop-down: Using private network endpoints to limit exposure and control access | Detecting prompt injection | Validating training data lineage | Setting retention policy

Answer: Using private network endpoints to limit exposure and control access

DO-08 Complete the Copilot Studio environment diagram: "DEV (Unmanaged) → TEST (Managed) → PROD (____)"

Drop-down: Managed | Unmanaged | Hybrid | Sandboxed

Answer: Managed

DO-09 Complete the AI Data ALM Gate table: "Gate B→C requires evidence of ____."

Drop-down: Profiling report; lineage IDs | CAB approval | Dashboards; alarms; budget guard | Residency mapping; approval memo

Answer: Profiling report; lineage IDs

DO-10 Complete the D365 Finance and Supply Chain environment table: "The environment used to validate with safe, anonymized production-like data is ____."

Drop-down: TEST | DEV | PROD | Sandbox-Preview

Answer: TEST

5. Case Studies (2 case studies, 8 questions total)

Case Study 1 — Contoso Financial Services

This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. To answer the case study questions, you will need to reference the information provided below. Information provided in an individual question does not apply to other questions in this case study.

Overview Contoso Financial Services is a mid-size lender that wants to deploy a Copilot Studio agent and a custom AI model to support loan-officer decision-making. The AI solution will read structured loan application data from Dynamics 365 Finance, unstructured supporting documents (PDF bank statements, ID scans), and telemetry from prior loan decisions.

Existing Environment

- Contoso currently has Dev, Test, and Production Power Platform environments.
- All financial data is classified as highly sensitive and subject to regional residency requirements.
- The security team requires managed identities for all AI workloads; no static credentials are permitted.
- The compliance team requires that no model be trained directly on live customer data.

Requirements

1. Loan officers must be able to ask the Copilot Studio agent to summarize an applicant's financial history.
2. A custom AI model must score default risk and must be evaluated on a golden dataset before every production release.
3. All prompts, tool calls, and safety events must be centrally logged and auditable.
4. Cross-region data movement must be blocked by default for this workload.

Case Study 1 — Question 1 You need to recommend an identity approach for the custom AI model's inference endpoint that satisfies the security team's requirement. What should you recommend?

- A. A shared service principal with a static secret rotated annually
- B. A managed identity assigned per environment (Dev, Test, Prod)
- C. The loan officer's own user credentials, delegated to the model
- D. An anonymous public endpoint secured only by network firewall rules

Answer: B

Case Study 1 — Question 2 You need to recommend where the default risk model should be trained, per the compliance requirement. What should you recommend?

- A. Train on live Production customer data, then anonymize the outputs
- B. Train/iterate using Dev/Test data only; never train on production knowledge
- C. Train on a public, third-party lending dataset with no internal review
- D. Train directly in Production during off-peak hours

Answer: B

Case Study 1 — Question 3 You need to recommend a logging approach that satisfies Requirement 3. What should you recommend?

- A. Per-agent local log files, reviewed manually each quarter
- B. Centralized logging that captures prompts, tool calls, errors, and safety events in a single workspace, with dashboards and alerting
- C. No logging for read-only summarization requests
- D. Logging only failed requests

Answer: B

Case Study 1 — Question 4 You need to recommend a setting that satisfies Requirement 4 (block cross-region data movement by default). What should you recommend?

- A. Set the default to in-region processing and require explicit approval to enable overflow processing
- B. Enable overflow processing globally to maximize model performance
- C. Leave the region setting unset so Copilot Studio decides automatically

- D. Disable residency controls for internal pilot workloads only

Answer: A

Case Study 2 — Northwind Manufacturing

This is a case study. To answer the case study questions, you will need to reference the information provided below. Information provided in an individual question does not apply to other questions in this case study.

Overview Northwind Manufacturing operates Dynamics 365 Supply Chain Management and wants to deploy an AI agent that predicts supply shortages and automatically drafts purchase orders for approval. The company also wants a repeatable ALM process for the agent, connectors, and underlying prediction model.

Existing Environment

- Northwind has Dev, Test, and Prod Power Platform environments with managed solutions used in Test and Prod only.
- The company uses Azure DevOps for release pipelines.
- Purchase order approval is considered a business-critical, high-risk action.
- The security team requires red team testing before any major release.

Requirements

1. The AI agent must never directly post a purchase order without a human approval step for actions that modify data.
2. Connector publishing to Prod must occur only after a full security review.
3. Before major releases, adversarial testing must occur and findings must be closed prior to go-live.
4. The prediction model must have a documented rollback plan before promotion to Prod.

Case Study 2 — Question 1 You need to recommend a control that satisfies Requirement 1. What should you recommend?

- A. Allow the agent to post purchase orders automatically to reduce latency
- B. Require approvals for publishing to production and for changes to high-risk capabilities, such as actions that modify data
- C. Allow any Maker to approve their own purchase order actions

- D. Remove the approval step for orders under a certain dollar threshold

Answer: B

Case Study 2 — Question 2 You need to recommend a process that satisfies Requirement 2. What should you recommend?

- A. Build connectors via Dev environment only, apply DLP during Test, and publish to Prod only after full security review
- B. Publish connectors directly to Prod from Dev to save time
- C. Allow any environment to publish connectors without review
- D. Skip DLP checks for internal-only connectors

Answer: A

Case Study 2 — Question 3 You need to recommend an action that satisfies Requirement 3. What should you recommend?

- A. Skip red teaming for incremental releases
- B. Run red team evaluations pre-production and after major changes; gate releases on findings closure
- C. Run red teaming only once per year regardless of release cadence
- D. Delegate red teaming entirely to the connector vendor

Answer: B

Case Study 2 — Question 4 You need to recommend a gate criterion that satisfies Requirement 4. What should you recommend?

- A. Promote to Prod once evaluation metrics exceed baseline, regardless of rollback documentation
- B. Confirm a rollback plan as part of Gate 2 (Test → Prod) criteria before promotion
- C. Document the rollback plan only after a Production incident occurs
- D. Rollback plans are only required for connectors, not models

Answer: B

6. Yes/No Hotspot Questions (3 statements per scenario; cannot be revisited after submission)

Hotspot Scenario 1 — Sales Executives Using Headset-Based Agent Queries

Sales executives use headsets to query an AI agent for account summaries during customer calls. The agent has access to CRM data and call transcripts.

For each statement, select **Yes** or **No**.

#	Statement	Answer	Reason
1	Should the AI agent escalate to a human after two failed attempts to resolve the executive's query?	Yes	Aligns with agent workflow governance requiring bounded tool-use and defined escalation paths rather than unbounded autonomous retries.
2	Should sensitive data be ignored (omitted) from logs by default?	No	Source requires audit logs to capture metadata for accountability; the correct control is redacting/sanitizing sensitive fields — not omitting logging altogether, since that removes auditability.
3	Should the agent be allowed to read call transcripts without any data classification applied?	No	Data governance requires sensitivity classification and DLP enforcement before an agent accesses any data source.

Hotspot Scenario 2 — Copilot Studio Agent Promotion

A Copilot Studio agent that summarizes customer cases is ready to move from Test to Production.

#	Statement	Answer	Reason
1	Can the release proceed if the connector has not completed a full security review?	No	"Publish connectors only in Prod after full review" is a stated ALM principle for connectors.
2	Is a documented rollback plan required before this promotion?	Yes	Rollback plans are a required gate artifact for production promotion across the ALM units.
3	Can a Maker publish this agent directly to Production without approval?	No	"No direct editing in production" and required publishing approvals apply.

Hotspot Scenario 3 — Custom Model Handling Financial Data

A custom AI model processes financial transaction data and produces anomaly alerts for the finance team.

#	Statement	Answer	Reason
1	Should the model be retrained directly on live Production financial data to stay current?	No	Explicit rule: never train on production knowledge; retrain using Dev/Test data, then re-evaluate before promotion.
2	Should prediction drift be monitored after the model goes live?	Yes	Explicitly required under Operational Monitoring for custom AI models.
3	Is it acceptable to skip the Model Card if the model passed all evaluation thresholds?	No	Model Card production is a required deliverable at the Evaluation & Approval stage, independent of whether thresholds were met.

Calibration Notes (carried over + new)

- DD-01 mirrors your deployment-method example almost exactly; note both agents land on the *same* answer pattern (automated pipeline) because the four stated requirements rule out manual methods for either agent equally — if you want differentiated answers, add a requirement that only applies to one agent type (e.g., "the Foundry agent must use infrastructure-as-code" vs. "the Copilot Studio agent must use managed solution packaging").
- DO-01 is now a direct restatement of your own example — verified against the source text on cross-app data flow testing.
- Case Study answers assume no partial credit; if your platform allows partial credit per sub-question, no changes needed.

Honest note on where I stretched beyond the source text: a few items (DD-09's tiered agent-persona access, DO-04's circuit-breaker/HITL pairing, and parts of Case Study 2 on threshold-based approval) extrapolate a stated principle into a new scenario rather than quoting a table verbatim — that's what **real 9/10 architect questions require, since**

the exam won't hand candidates the answer in a table either. But if a candidate pushes back that "the doc doesn't say this exact combination," they'd have a legitimate point on those specific items, so I'd flag them for a second SME pass before they go live in a scored exam rather than treating them as bulletproof like the gate-table questions are.

AB-100 Exam Question Bank — v3 (Architect Level, Difficulty 9/10)

These questions require synthesizing multiple units at once (ALM gates + security + residency + testing + governance in a single scenario), the way a real solution architect has to reconcile competing constraints for one customer. Every "wrong" option is wrong for a specific, defensible reason — not just a distractor.

1. Drag-and-Drop (10) — Architect Level

DD-01 — Multi-Agent Deployment Under Conflicting Constraints A global insurer is deploying three AI components for a claims-automation solution:

- A Copilot Studio agent that drafts customer correspondence (low risk, high change frequency)
- A Microsoft Foundry agent that adjudicates claims under \$5,000 autonomously (high risk, regulated action)
- A custom fraud-detection model retrained monthly (high risk, drift-sensitive)

Constraints:

- The Foundry agent's autonomous adjudication is a "data-modifying, high-risk action" requiring the strictest promotion controls.
- The fraud model must never see production claims data during training.
- The correspondence agent changes weekly and cannot tolerate multi-day release cycles.

Match each component to the correct promotion/deployment pattern.

Items: Fast automated pipeline with lightweight Gate 1 checks and frequent Dev→Test→Prod cycles | Full Gate 1 + Gate 2 promotion with mandatory human validation of reasoning, rollback plan, and CAB approval before Prod | Automated pipeline with strict Dev/Test-only training data enforcement, golden-set evaluation, and Model Card produced before promotion

Targets: Copilot Studio correspondence agent | Foundry claims-adjudication agent | Fraud-detection model

Answer:

- Correspondence agent → Fast automated pipeline, lightweight Gate 1, frequent cycles (low risk, high change velocity — heavyweight CAB gating here would be over-engineering and fail the "cannot tolerate multi-day cycles" constraint)
- Foundry adjudication agent → Full Gate 1 + Gate 2 with human validation, rollback, CAB approval (high-risk, data-modifying action explicitly requires this per Promotion Gate 2 criteria)
- Fraud-detection model → Automated pipeline with Dev/Test-only training enforcement + golden-set evaluation + Model Card (matches the "never train on production data" constraint and custom-model ALM requirements)

Architect trap: a naive answer applies the same heavyweight gate to all three, which fails the correspondence agent's velocity constraint and is over-governed; an equally naive answer applies the lightweight pattern to all three, which violates the adjudication agent's regulatory risk.

DD-02 — Residency vs. Latency Trade-off A EU-based bank requires in-region data processing by default. A US subsidiary needs low-latency access to the same Copilot Studio agent for a pilot. Order the correct decision sequence per the Residency Decision Tree, given the pilot is *not* pre-approved for overflow.

Items: Block the feature for the US subsidiary or defer until overflow is approved | Check if in-region capacity is available | Check if overflow is allowed for this workload tier | Enable cross-region processing under admin control

Answer (correct order): Check if in-region capacity is available → (No) → Check if overflow is allowed for this workload tier → (Not pre-approved = No) → Block the feature for the US subsidiary or defer until overflow is approved (*Note: "Enable cross-region processing under admin control" is a distractor — it only applies if overflow IS allowed, which this scenario explicitly rules out.*)

DD-03 — Layered Security Control Mapping Under a Prompt-Injection Incident A support agent was found to leak internal pricing data after a crafted prompt-injection attempt embedded in an uploaded PDF. Match each remediation to the correct control layer it primarily addresses.

Items: Sanitize tool inputs and strip unsafe embedded content from uploaded files | Apply DLP policy to block pricing data from appearing in outputs | Restrict the agent's connector scope so pricing tables aren't queryable at all | Add anomaly alerting for unusual data-access spikes

Targets: Input/output filtering | Data governance (DLP/sensitivity) | Identity & access (least privilege) | Observability

Answer:

- Input/output filtering → Sanitize tool inputs and strip unsafe embedded content
- Data governance → Apply DLP policy to block pricing data in outputs
- Identity & access → Restrict connector scope (least privilege)
- Observability → Anomaly alerting for data-access spikes (*Architect insight: this is a layered-defense question — a single fix is insufficient; the real answer is that all four apply simultaneously, but each maps to a distinct control layer, which is what a defense-in-depth review actually checks for.*)

DD-04 — Ordering a Cross-Module ALM Release for a Regulated Workload A

pharmaceutical company's Foundry agent uses grounding data from Dataverse and a custom model trained on clinical trial data. Order the correct end-to-end release sequence combining data ALM and agent ALM.

Items: Freeze gold datasets and sign immutability attestations | Deploy validated agent bundle to Prod with region/residency enforcement | Run evaluation suite on golden sets and produce Model Card | Execute canary runs using masked/representative Prod-like data | Ingest and profile clinical trial data, generate lineage

Answer (correct order): Ingest and profile clinical trial data, generate lineage → Run evaluation suite on golden sets and produce Model Card → Execute canary runs using masked/representative Prod-like data → Freeze gold datasets and sign immutability attestations → Deploy validated agent bundle to Prod with region/residency enforcement

DD-05 — Matching Testing Types to a Multi-App Regulatory Scenario A logistics firm needs to validate an AI solution spanning Dynamics 365 Supply Chain, Finance, and a Foundry agent that reroutes shipments automatically. Match the correct testing type to each requirement.

Items: Confirms shipment reroute logic doesn't violate customs/compliance rules under edge-case inputs | Confirms data consistency when a shipment update in Supply Chain triggers a Finance ledger entry | Confirms the agent doesn't regress after a prompt template change | Confirms the reasoning quality of the agent's autonomous rerouting decision

Targets: Scenario/edge-case testing | Integration/end-to-end testing | Regression testing | Model/agent behavior evaluation

Answer:

- Customs/compliance edge cases → Scenario/edge-case testing
 - Supply Chain→Finance data consistency → Integration/end-to-end testing
 - No regression after prompt change → Regression testing
 - Reasoning quality of autonomous decision → Model/agent behavior evaluation
-

DD-06 — Order the Incident Response Sequence for a Compromised Model

Endpoint A model endpoint is suspected of leaking training data. Order the correct incident response sequence.

Items: Restore service using a prior validated version after root cause is addressed | Disable the compromised endpoint | Notify stakeholders and begin investigation | Preserve inference logs and model artifacts for forensic analysis

Answer (correct order): Disable the compromised endpoint → Preserve inference logs and model artifacts for forensic analysis → Notify stakeholders and begin investigation → Restore service using a prior validated version after root cause is addressed

DD-07 — Matching RACI Roles to a Cross-Border Data Movement Decision A multinational retailer needs sign-off to enable cross-region overflow for an AI workload. Match each role to its RACI designation for "Residency configuration," per the AI data RACI table.

Items: Data Owner | AI Architect | Security/Compliance | Platform Admin

Targets: Consulted | Accountable | Responsible | Responsible

Answer: Data Owner → Consulted; AI Architect → Accountable; Security/Compliance → Responsible; Platform Admin → Responsible (*Architect trap: many candidates assume Security/Compliance is Accountable for a compliance decision — but the source table places Accountability with the AI Architect, who owns the overall design decision, while Security/Compliance is Responsible for execution.*)

DD-08 — Order a Model Retraining Cycle Triggered by Detected Drift A demand-forecasting model shows accuracy degradation flagged by drift monitoring. Order the correct remediation sequence.

Items: Redeploy through Gate 1 and Gate 2 promotion checks | Detect drift vs. baseline metrics | Retrain using updated, versioned Dev/Test data | Re-run evaluation suite against golden sets | Trigger an incident/backlog item with trace IDs

Answer (correct order): Detect drift vs. baseline metrics → Trigger an incident/backlog item with trace IDs → Retrain using updated, versioned Dev/Test data → Re-run evaluation suite against golden sets → Redeploy through Gate 1 and Gate 2 promotion checks

DD-09 — Matching Governance Controls to Agent Personas in a Zero-Trust Design A manufacturer is designing three agent personas: a read-only analytics agent, an agent that updates inventory records, and an agent that can trigger supplier payments. Match the appropriate access pattern to each persona.

Items: Read-only RBAC scope with no write actions enabled | Scoped write access to inventory tables only, with DLP on export actions | Scoped write access plus mandatory multi-step human approval before execution, with full audit trail

Targets: Analytics agent | Inventory-update agent | Supplier-payment agent

Answer:

- Analytics agent → Read-only RBAC scope
 - Inventory-update agent → Scoped write access to inventory tables + DLP on export
 - Supplier-payment agent → Scoped write access + mandatory multi-step approval + full audit trail (*This mirrors "Restrict privileged access... require multi-step approvals for sensitive model updates" applied to agent action design, not just model updates — a legitimate extrapolation an architect must reason through.*)
-

DD-10 — Order the Responsible AI Review Combined with a Fairness Finding During a Responsible AI review, a fairness testing pass reveals disparate outcomes for one customer segment. Order the correct remediation sequence.

Items: Redeploy after governance board sign-off | Document the fairness gap and root cause in the RAI review report | Apply mitigation (data rebalancing, threshold adjustment) and re-test | Escalate to governance board for review of significant model changes | Detect and validate the disparate impact using fairness testing on real-world scenarios

Answer (correct order): Detect and validate the disparate impact using fairness testing → Document the fairness gap and root cause in the RAI review report → Apply mitigation

and re-test → Escalate to governance board for review of significant model changes → Redeploy after governance board sign-off

2. Drop-Down / Diagram Completion (10) — Architect Level

DO-01 A bank's AI architecture diagram shows: Data Sources → Agent Runtime → [___] → User Output & Logging. The bank has just added a *new* requirement: block outbound transmission of any field carrying a "Highly Confidential" sensitivity label, even if DLP rules haven't been explicitly written for that field yet. Which layer, and what capability within it, satisfies this?

Drop-down: DLP & Sensitivity Filters, configured to surface/enforce the highest sensitivity label present, not just explicit DLP rules | Agent Runtime, with token limits reduced | User Output & Logging, with output truncated | Data Sources, with the field deleted from the source system

Answer: DLP & Sensitivity Filters, configured to surface/enforce the highest sensitivity label present — matches "Apply sensitivity labels to knowledge sources; surface the highest label in responses where supported," which is broader than relying on pre-written DLP rules alone.

DO-02 An architect is completing a Gate table for a Foundry agent that also touches Dynamics 365 Finance data. The customer asks: "At which gate do we confirm the agent's cost and performance under real ERP workload constraints, not synthetic load?"

Drop-down: Gate 2 (Test → Prod) — Performance verified under ERP workload constraints | Gate 1 (Dev → Test) — Initial safety checks | Gate A→B — Catalog & ownership | Gate D→E — Compliance & residency

Answer: Gate 2 (Test → Prod) — Performance verified under ERP workload constraints (D365 F&SC ALM Gate 2 criteria, distinct from the generic AI-data Gate D→E which is about compliance/residency, not ERP load).

DO-03 Complete the Model Hardening Blueprint for a scenario where the customer specifically flags *training-time* poisoning as their top concern (not runtime attacks). Which step of the blueprint should be emphasized most, and why?

Drop-down: Step 4, Validation Pipeline — validate models and data to maintain quality and compliance, including new-data validation before retraining | Step 2, Private Endpoints — network isolation | Step 3, Threat Protection — runtime detection | Step 1, Secure Compute — compute isolation

Answer: Step 4, Validation Pipeline — poisoning is a training-data integrity risk; the blueprint's guidance to "validate newly introduced data before using it in retraining pipelines" is the direct control, not network isolation or runtime threat detection (which address different attack surfaces).

DO-04 A customer wants their architecture diagram to show where "Human-in-the-loop" (HITL) sits relative to drift detection for a fraud model. Complete: "Detect data drift vs. baselines → trigger safeguard actions (____, HITL)."

Drop-down: Circuit breakers | Load balancers | CDN caching | Autoscaling rules

Answer: Circuit breakers — explicitly paired with HITL in the source: "Detect data drift vs. baselines; trigger safeguard actions (circuit breakers, HITL)."

DO-05 Complete the RBAC design diagram where the customer wants to know which role should hold approval authority for publishing to Production and for high-risk capability changes.

Drop-down: Security Admin / Publisher approval workflow, distinct from Maker | Maker role, self-approved | Environment Admin only, no separate approval needed | Any authenticated user with a Copilot Studio license

Answer: Security Admin / Publisher approval workflow, distinct from Maker — "Distinct roles for Maker, Publisher, Environment Admin, and Security Admin... Require approvals for publishing to production."

DO-06 A regulated customer's Foundry agent diagram must show where secrets are separated from development artifacts. Complete: "Data, security, and residency controls" must enforce ____.

Drop-down: Separation of production secrets from development artifacts, using secure identity management for agent actions | Shared secrets vault across all environments for simplicity | Secrets embedded directly in agent prompt templates | Secrets stored in the source code repository for traceability

Answer: Separation of production secrets from development artifacts, using secure identity management for agent actions

DO-07 Complete the audit-trail architecture diagram: the customer wants Foundry activity logs exported to their existing SIEM. Which two Foundry-native capabilities does this diagram need?

Drop-down: Foundry activity logs exported to Azure Monitor / Log Analytics / SIEM tools (e.g., Microsoft Sentinel), plus diagnostics/tracing for execution path visibility | A manually maintained spreadsheet of changes | A single unstructured text log file per model | Screenshots of the model registry taken weekly

Answer: Foundry activity logs exported to Azure Monitor / Log Analytics / SIEM tools, plus diagnostics/tracing

DO-08 A customer's compliance officer asks the architect to complete this statement for their audit-retention diagram: "Regulated workloads should retain audit logs for ____."

Drop-down: 12–24 months, per the recommended pattern for regulated workloads (vs. 90 days for low-risk, indefinite for incident-related archives) | Exactly 30 days for all workloads | Indefinitely for all workloads regardless of risk | No retention requirement if logs contain no PII

Answer: 12–24 months, per the recommended pattern for regulated workloads

DO-09 Complete the grounding-data access diagram for a scenario where the customer wants model-tuning operations completely isolated from production inference. What environment design principle fills the blank: "Model tuning access requires ____."

Drop-down: Segregated environments for development, evaluation, and production, insulated from production operations | The same environment as production inference, with extra logging | A shared identity between tuning and inference workloads | No environment segregation, since tuning is a one-time event

Answer: Segregated environments for development, evaluation, and production, insulated from production operations

DO-10 A customer's architecture review flags that their agent's connector allow-list is undocumented. Complete the Development and Interoperability Standards diagram: "Agents must use ____ for structured tool/data access, and ____ for controlled delegation across agents."

Drop-down: Model Context Protocol (MCP); Agent-to-Agent (A2A) | REST APIs only; SOAP for legacy systems | GraphQL; Webhooks | OData; SFTP

Answer: Model Context Protocol (MCP); Agent-to-Agent (A2A) — explicitly named as the standard protocols for tool/data access and controlled agent delegation.

3. Case Studies (2–3 scenarios, 8 questions total) — Architect Level

Case Study 1 — Meridian Health Network (Regulated, Multi-Region)

This is a case study. Information provided in an individual question does not apply to other questions in this case study.

Overview Meridian Health Network operates hospitals across the EU and US. It is deploying a Foundry agent that summarizes patient intake notes for clinicians and a custom AI model that flags high-risk readmission candidates. The solution draws on Dataverse patient records, a Dynamics 365 Customer Service case history, and a third-party lab-results API.

Existing Environment

- EU patient data must remain in-region by regulatory mandate; no pre-approved overflow exists for this workload tier.
- The readmission model is retrained quarterly using updated clinical data.
- The intake-summarization agent is considered a high-risk, data-sensitive capability because it surfaces PHI (protected health information).
- Meridian's security team mandates managed identities for every agent and pipeline, with no embedded secrets anywhere.
- A prior incident involved a prompt-injection attempt via a malicious PDF lab report; leadership now requires this class of attack be mitigated architecture-wide, not just patched once.

Requirements

1. EU clinician queries must never trigger cross-region processing.
2. The readmission model must never train on live PHI directly from Production.
3. The intake agent must have layered protections against embedded malicious content in uploaded documents, not a single point fix.
4. Every promotion of the readmission model to Production must include a documented rollback plan and CAB approval.

Case Study 1 — Q1 A US-based data scientist proposes routing EU clinician queries through a US region temporarily during a capacity shortage, without new approvals,

citing the pilot exemption used in an unrelated non-clinical workload. What should the architect do?

- A. Approve the routing since a pilot exemption previously existed for a different workload
- B. Reject the routing; per the Residency Decision Tree, overflow requires explicit approval for *this* workload tier, and none exists — in-region capacity issues should block or defer the feature, not silently reroute
- C. Approve the routing only for read-only queries
- D. Approve the routing but disable logging to reduce audit exposure

Answer: B — Precedent from an unrelated workload does not satisfy the explicit, per-workload-tier overflow approval requirement; this is a classic architect trap where surface-level precedent is used to justify a policy violation.

Case Study 1 — Q2 Which data preparation approach satisfies Requirement 2 while still keeping the quarterly retraining realistic and useful?

- A. Train directly on live Production PHI but delete the training copy after each run
- B. Use versioned, de-identified/curated Dev/Test datasets refreshed each quarter from a controlled pipeline, with lineage and hashes stamped per snapshot
- C. Train on the third-party lab API's public sample dataset only, ignoring internal clinical data entirely
- D. Skip retraining and rely on the original launch model indefinitely

Answer: B — Matches "never train on production knowledge" combined with the versioned, lineage-stamped curated dataset approach from the AI Data ALM process (Phase B — Ingest & Prepare).

Case Study 1 — Q3 Given the prior prompt-injection incident and the requirement for layered (not single-point) protection, which combination of controls should the architect recommend for the intake agent?

- A. Input sanitization only — strip unsafe embedded content from uploaded PDFs
- B. DLP policy only — prevent PHI from appearing in outputs
- C. A layered stack: input/output filtering (sanitize uploaded content), data governance (DLP/sensitivity labels on PHI), least-privilege connector scope, and

observability (anomaly alerting on unusual access patterns) — combined, not standalone

- D. Rely entirely on the underlying model's built-in safety training, since it already resisted the original attack

Answer: C — This is the defense-in-depth principle applied directly; any single control (A, B, or D alone) leaves other attack surfaces exposed and fails "mitigated architecture-wide, not just patched once."

Case Study 1 — Q4 Which gate requirement, if skipped, would most directly violate Requirement 4?

- A. Skipping the CAB approval step because the evaluation metrics exceeded baseline
- B. Skipping the initial Dev-stage prompt drafting review
- C. Skipping the data profiling step during ingestion
- D. Skipping the connector authentication validation step

Answer: A — Requirement 4 explicitly ties CAB approval and rollback plan documentation to every Production promotion, independent of how strong the evaluation metrics are; this is the same trap as Gate 2's rollback-plan requirement being metric-independent.

Case Study 2 — Vantage Global Logistics (Multi-App, High Automation)

This is a case study. Information provided in an individual question does not apply to other questions in this case study.

Overview Vantage Global Logistics runs Dynamics 365 Supply Chain Management and Dynamics 365 Finance. It is deploying a Foundry agent that autonomously reroutes shipments and automatically posts related cost adjustments to the Finance ledger when reroutes occur. This is explicitly flagged as a business-critical, data-modifying automation.

Existing Environment

- The company uses Azure DevOps for all release pipelines; managed solutions are used in Test and Prod only.
- A recent audit found inconsistent testing: some releases skipped end-to-end validation of the Supply Chain → Finance data handoff.

- The agent's connector to the Finance ledger was flagged as high-risk because it can post financial entries.
- Leadership wants confidence that if the agent's reasoning is later found flawed, the team can trace exactly which release and data snapshot caused it.

Requirements

1. The agent must not autonomously post ledger entries above a materiality threshold without human approval.
2. Every release must validate that a Supply Chain reroute event correctly and consistently produces the matching Finance ledger entry.
3. The connector used to post to Finance must go through full security review before every Production publish, not just the first time.
4. The team must be able to trace any flawed ledger entry back to a specific model/agent version and specific data snapshot.

Case Study 2 — Q1 Which access-pattern design satisfies Requirement 1 without eliminating the automation's core value?

- A. Full autonomous posting with no human step, since the agent passed initial testing
- B. Scoped write access to post ledger entries, with mandatory human approval required only above a defined materiality threshold, and full audit trail on every action
- C. Read-only access — the agent may only recommend but never post anything
- D. Remove the agent's Finance connector entirely and route everything through email approval

Answer: B — This mirrors the tiered persona-access design pattern (from DD-09): full autonomy is too risky for a data-modifying, business-critical action, but pure read-only defeats the automation's purpose; the threshold-based approval is the correct middle design.

Case Study 2 — Q2 Which testing type directly addresses the audit finding referenced in Requirement 2?

- A. Unit testing of the reroute algorithm in isolation
- B. Integration/end-to-end testing validating cross-app data flow between Supply Chain and Finance

- C. UI regression testing of the Supply Chain dashboard
- D. Load testing of the Finance API under peak traffic

Answer: B — Directly matches the source's rationale for why end-to-end tests must validate cross-app data flow: "AI output quality depends on consistent, trusted, and well-timed input data from across integrated systems."

Case Study 2 — Q3 Which ALM principle, applied correctly, satisfies Requirement 3 (repeated security review, not one-time)?

- A. Publish connectors to Prod once security review passes at initial launch, then treat subsequent releases as pre-approved
- B. Apply the full connector release flow (Author → Validate Auth → Apply DLP → Approve Security → Publish to Prod) on every release that touches this connector, not only the first
- C. Delegate all future security reviews to the Platform Admin without Security team involvement
- D. Skip DLP re-validation if the connector's authentication method hasn't changed

Answer: B — The connector ALM release flow is a per-release gate, not a one-time approval; treating it as "grandfathered" after the first pass is the exact trap this question tests.

Case Study 2 — Q4 Which combination of controls best satisfies Requirement 4 (traceability of a flawed ledger entry to a specific version and data snapshot)?

- A. Immutable, timestamped, role-attributed audit logs for model/agent version changes, correlated with versioned/lineage-stamped data snapshots and tracing fields (correlation ID, model version)
- B. A weekly manual export of the model registry to a spreadsheet
- C. Reliance on the Finance team's own ledger reconciliation process alone
- D. Deleting older model versions after each promotion to reduce storage cost

Answer: A — Combines audit-trail requirements (immutable logs, role attribution) with data lineage/versioning and tracing fields (correlation ID, model version) — this is the only option that actually enables the traceability leadership asked for; D is actively counterproductive since it destroys the very versions needed for tracing.

4. Yes/No Hotspot Questions — Architect Level (3 statements per scenario; cannot be revisited after submission)

Hotspot Scenario 1 — Sales Executives Using Headset-Based Agent Queries (Extended, Multi-Constraint)

Sales executives use headsets to query an AI agent for account summaries during live customer calls. The agent has access to CRM data, call transcripts, and a connector to a pricing system. The company has a strict "no direct production edits" policy and requires that any escalation path preserve context so a human doesn't have to ask the customer to repeat themselves.

#	Statement	Answer	Reason
1	Should the agent be allowed to modify pricing directly through its connector without a separate approval step, given it's only "assisting" a live call?	No	Pricing changes are a data-modifying, business-critical action; per governance principles, high-risk capability changes require approval workflows regardless of the "assist-only" framing — the risk is in the connector's write capability, not the agent's stated intent.
2	If the agent escalates to a human after repeated failures, should the escalation preserve the conversation context (transcript, CRM lookups already performed) rather than starting cold?	Yes	This follows directly from the requirement that a human "doesn't have to ask the customer to repeat themselves" — architecturally this means the escalation path must carry forward session context, consistent with the observability/tracing guidance (correlation IDs, session continuity) rather than treating escalation as a hard reset.
3	Is it acceptable to log the full pricing figures discussed on the call in plaintext for training-data purposes, since pricing isn't classified as PII?	No	Sensitivity classification isn't limited to PII — pricing data can carry its own sensitivity label per DLP policy design ("Apply sensitivity labels to knowledge sources; surface the highest label"), and using live call content for training without going through the curated Dev/Test pipeline violates "never train on production knowledge."

Hotspot Scenario 2 — Cross-Border M&A Data Consolidation

A multinational is merging two subsidiaries' Dynamics 365 environments, each in a different data residency region. An AI agent will be used temporarily to reconcile duplicate customer records across both regions during the migration window.

#	Statement	Answer	Reason
1	Is it acceptable to run this reconciliation agent with a single shared identity across both regional environments to simplify the temporary migration?	No	Even for temporary/migration workloads, the principle "assign a unique cloud identity per agent, per environment... with recorded ownership, version, and lifecycle metadata" still applies; a shared identity removes traceability exactly when cross-region data handling needs it most.
2	Does cross-region record matching during this migration require the same in-region-first, explicit-overflow-approval decision process as any other AI workload?	Yes	The Residency Decision Tree doesn't carve out an exception for temporary/migration workloads — "temporary" is not one of the stated exemption criteria; the same explicit approval gate applies.
3	Once the migration is complete, is it acceptable to leave the reconciliation agent's elevated cross-region access permissions active indefinitely "in case it's needed again"?	No	This directly contradicts least-privilege and lifecycle governance principles ("criteria for archiving or retiring unused agents"); temporary elevated access must be retired, not left standing as a convenience.

Hotspot Scenario 3 — Autonomous Supplier Payment Agent Under Drift

An AI agent autonomously approves supplier payments under a defined materiality threshold. Six months after go-live, drift monitoring shows the agent's approval pattern has shifted meaningfully from its original validated baseline, though no individual payment has yet caused a loss.

#	Statement	Answer	Reason
1	Because no financial loss has occurred yet, is it acceptable to defer	No	Drift detection guidance requires that when drift exceeds thresholds, the system should "auto-open an incident,

	investigation until the next scheduled quarterly review?		route to data owner, and pause affected actions" — the trigger is the drift signal itself, not realized loss; waiting for harm before acting inverts the intended safeguard.
2	Should the agent's autonomous approval capability be paused (or its threshold tightened) while the drift is investigated?	Yes	This is the explicit safeguard action tied to drift detection — "trigger safeguard actions (circuit breakers, HITL)" — pausing or narrowing scope is the circuit-breaker response, not a discretionary choice.
3	If the root cause is found to be a legitimate shift in supplier behavior (not a model defect), can the team simply raise the drift-monitoring threshold to stop the alerts, without retraining or re-evaluating the model?	No	The correct response to confirmed drift — even a "legitimate" shift — is to retrain on updated, versioned data and re-run evaluation against golden sets before redeployment, per the retrain/re-evaluate/redeploy cycle; silencing the alert by loosening thresholds bypasses governance rather than resolving the underlying model-data mismatch.

Architect-Level Calibration Notes

- Every "wrong" option above is wrong for a *specific, named* reason traceable to source guidance — not generic distractor noise. This is what separates a 9/10 difficulty item from a recall item: the incorrect options represent plausible real-world shortcuts (precedent reuse, "it's just a pilot," "no loss yet") that an under-experienced architect might actually take.
- Case Study 2's Q4 (traceability) and Hotspot Scenario 3 both hinge on the same underlying principle (drift/incident response is trigger-based, not harm-based) tested from two different angles — intentional, to check whether the candidate generalizes the principle or only memorized one phrasing of it.
- If your exam platform enforces "no revisit" per hotspot block, double check Hotspot Scenario 1, Q1 vs Q3: a candidate who gets Q1 wrong (allows autonomous pricing edits) is likely to also misjudge Q3 in the same direction — this is intentional correlation to catch candidates who don't understand the underlying least-privilege principle versus those who got lucky on one item.